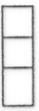


## Introduction

Data in most large institutions are managed by heterogeneous and autonomous application systems and their databases. It is well understood that data is an institutional resource. However, supporting the desired level of interoperability, uniform access, and/or consistency of data will likely remain technically and organizationally a challenging task.

In this tutorial, we will study the issues and the techniques that can support interoperability in multidatabase systems from both, database and application perspectives.



# ARCHITECTURE

## SYSTEM ARCHITECTURE

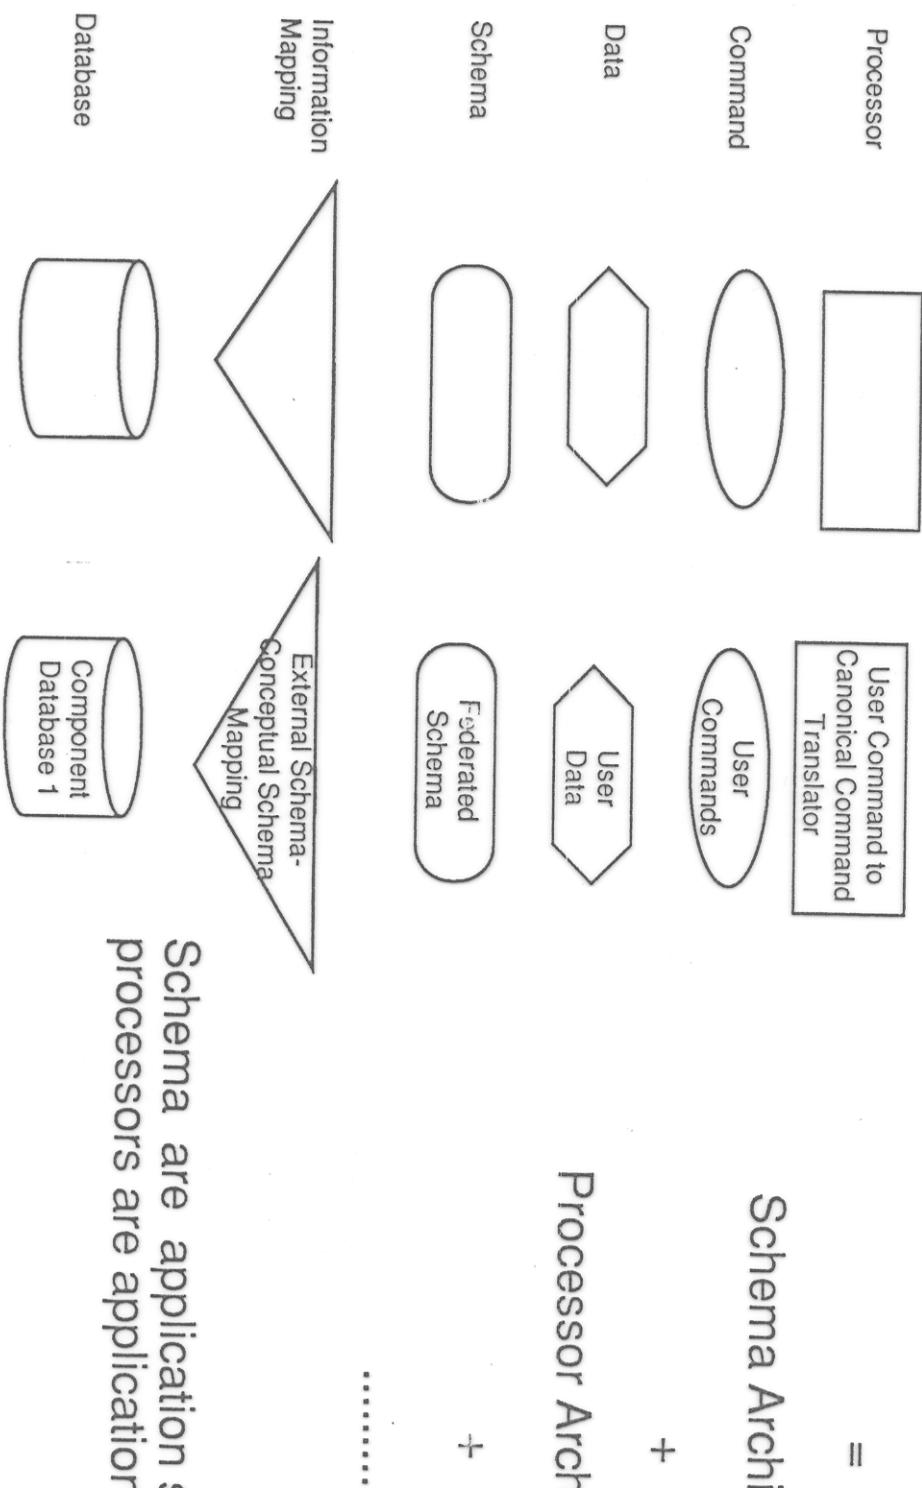
Component Type	Icon	Example
Processor		User Command to Canonical Command Translator
Command		User Commands
Data		User Data
Schema		Federated Schema

## Processor Architecture

+

## Schema Architecture

+



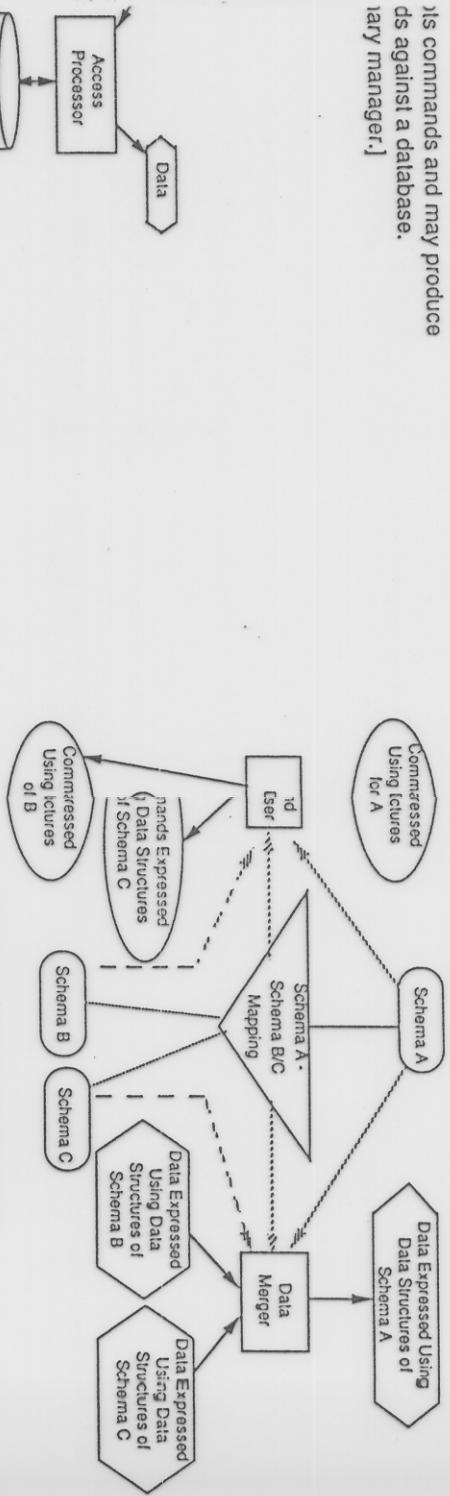
Schema are application specific,  
processors are application independent

[Sheth and Larson 90]

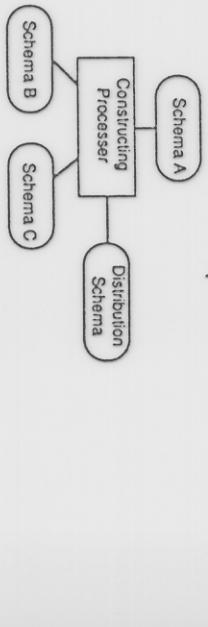
## Constructing Processors

liting processors partition and replicate operations produced by a processor into operations that are accepted by two or more processors. They also merge data produced by several processors. They support location, distribution and replication transparencies, query decomposer and optimizer, global transaction manager.]

] Processor  
Is commands and may produce  
ds against a database.  
ary manager.]



(a) A pair of Constructing Processors.

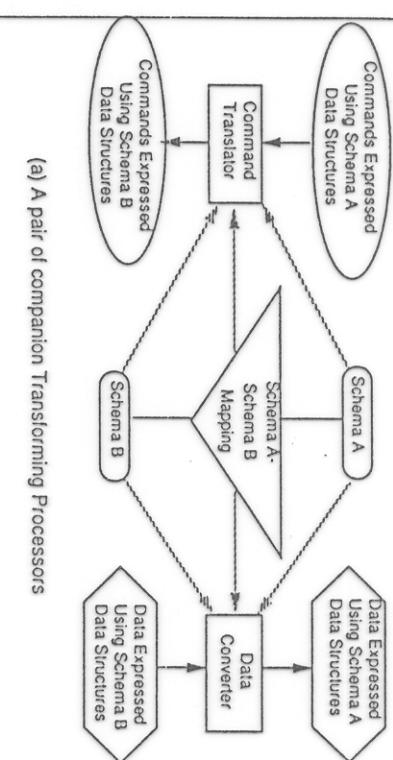


(b) An Abstract Constructing Processor.

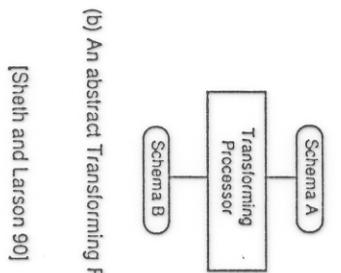
[Sheh and Larson 90]

## Transforming Processors

Transforming processors translate commands (or operations) from one language or format to another language or format, or transform data from one format to another. They can be used to support data model transparency. [E.g., A SQL- CODASYL command transformer.]



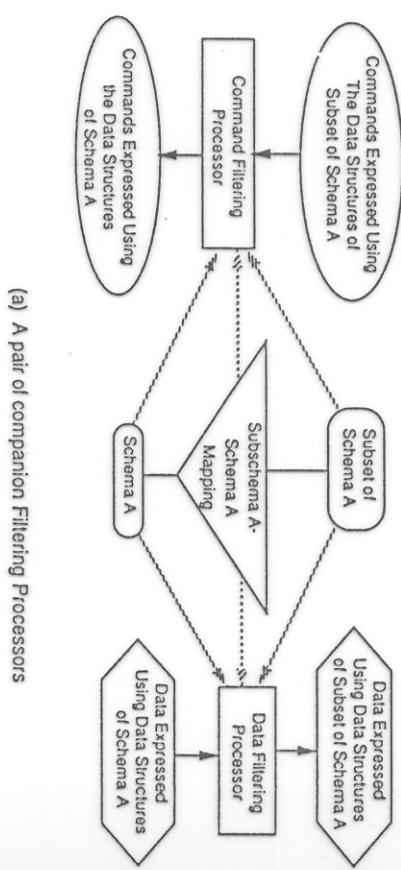
(a) A pair of companion Transforming Processors



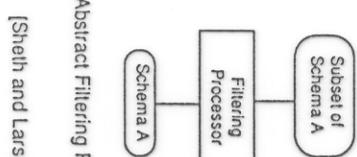
(b) An abstract Transforming Processor

## Filtering Processors

Filtering processors constrain operations that can be applied, or constrains data that can be passed to another processor. This gives ability to control access of one processor to another. [E.g., access controller.]



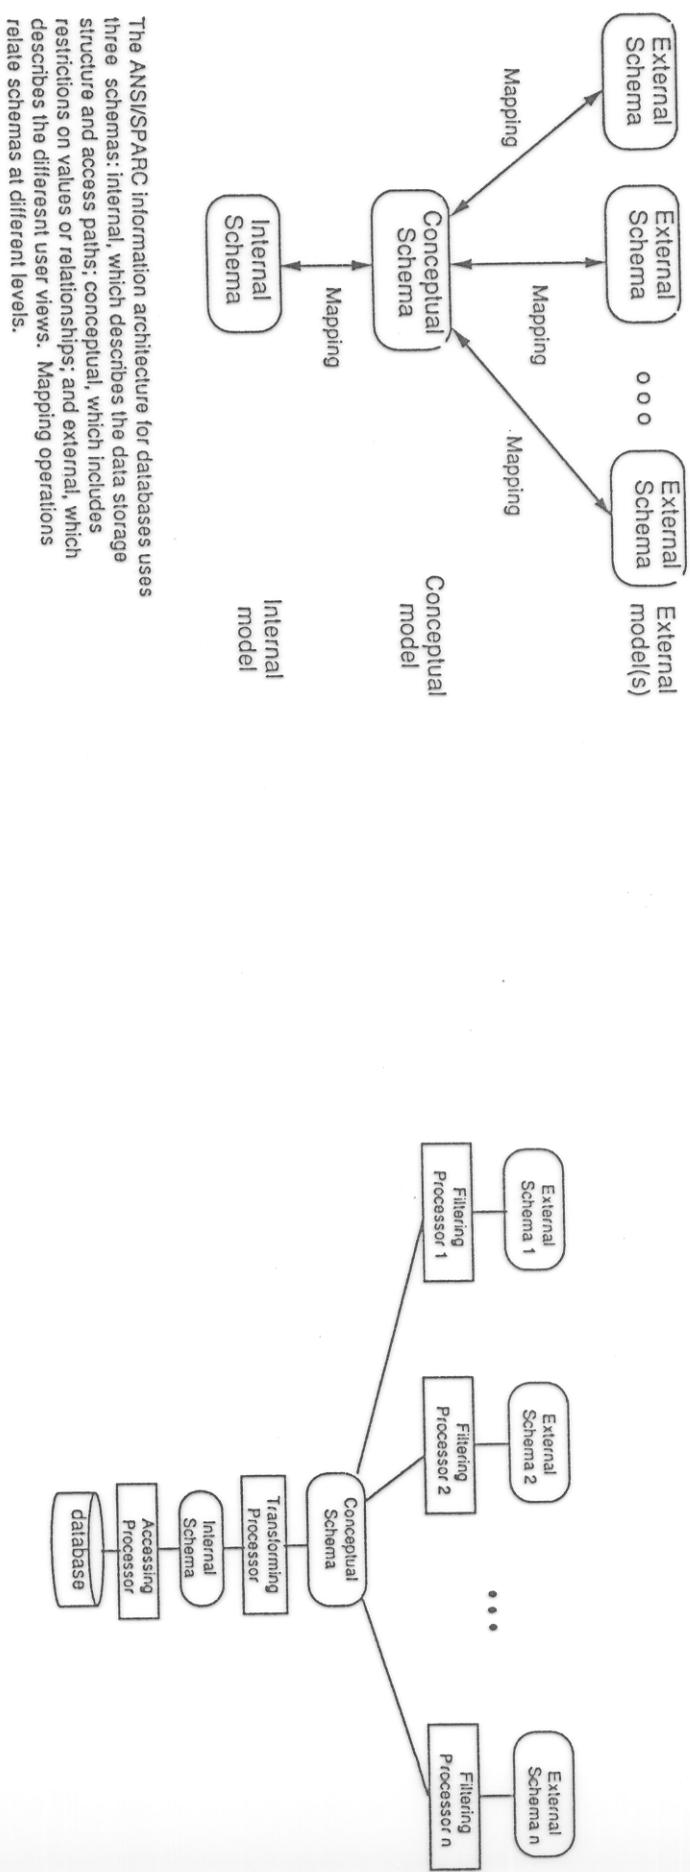
(a) A pair of companion Filtering Processors



(b) An Abstract Filtering Processor.

## Three-Level Schema Architecture of a Centralized DBMS

### System Architecture of a Centralized DBMS



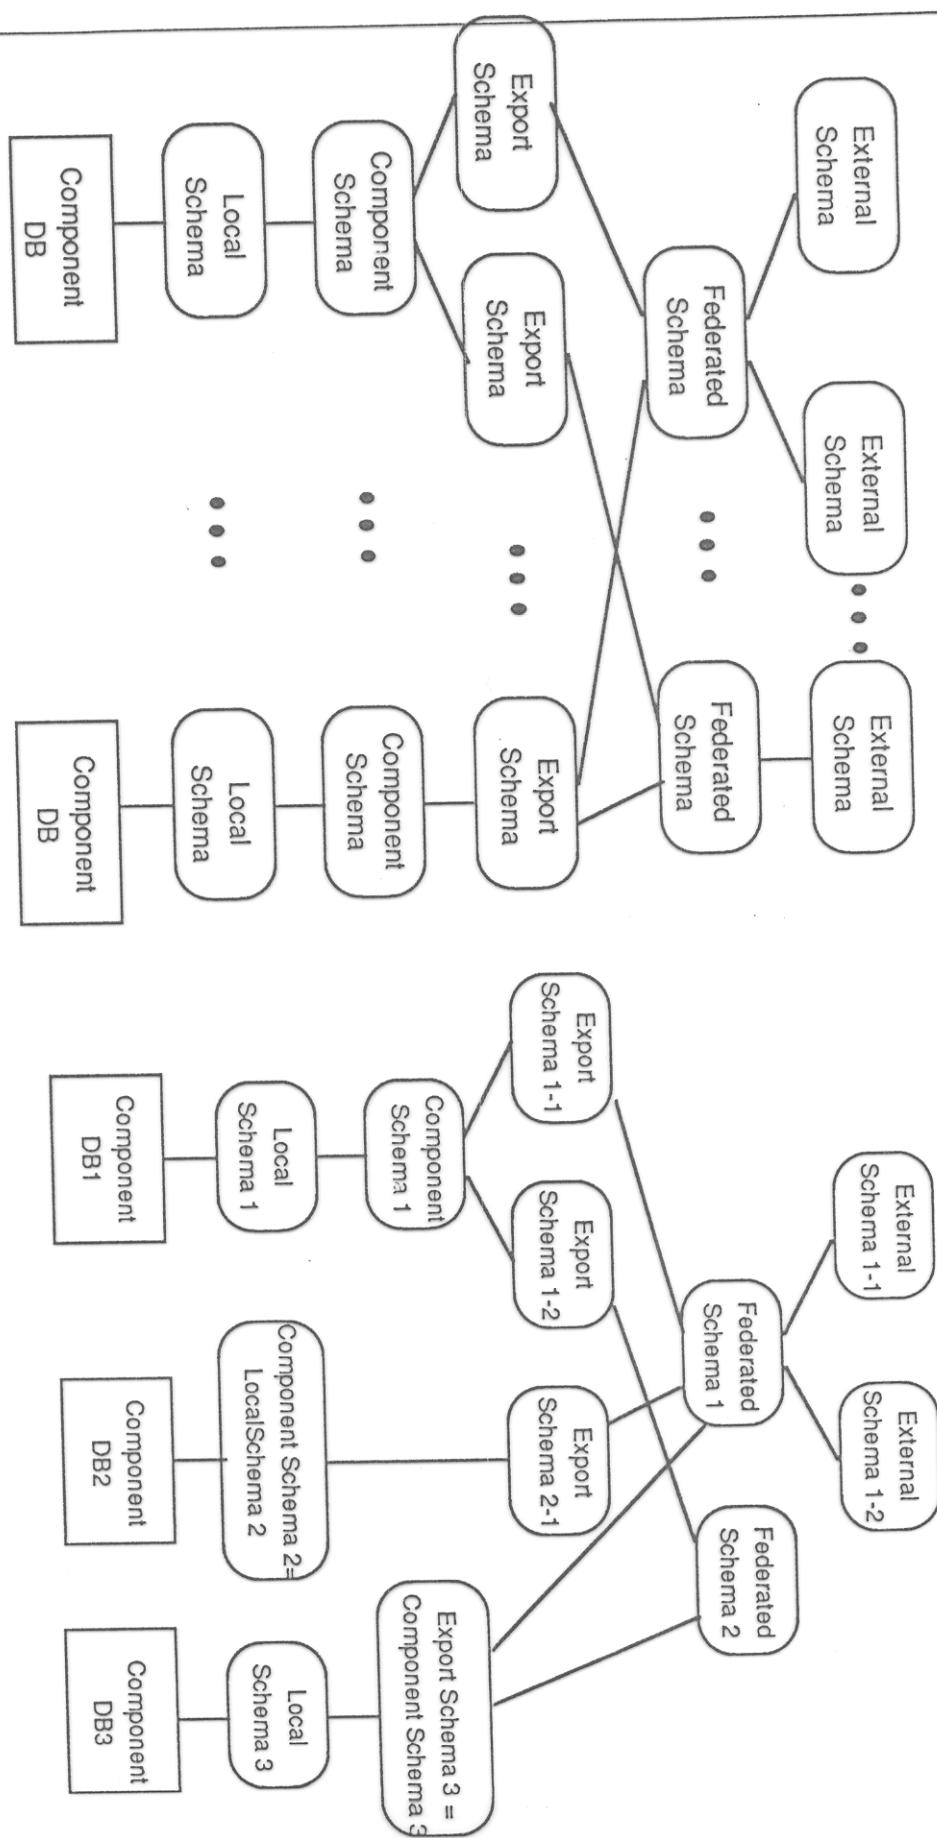
The ANSI/SPARC information architecture for databases uses three schemas: internal, which describes the data storage structure and access paths; conceptual, which includes restrictions on values or relationships; and external, which describes the different user views. Mapping operations relate schemas at different levels.

[ANSI/SPARC 77]

[Sheith and Larson 90]

## Five-Level Schema Architecture of a FDBS

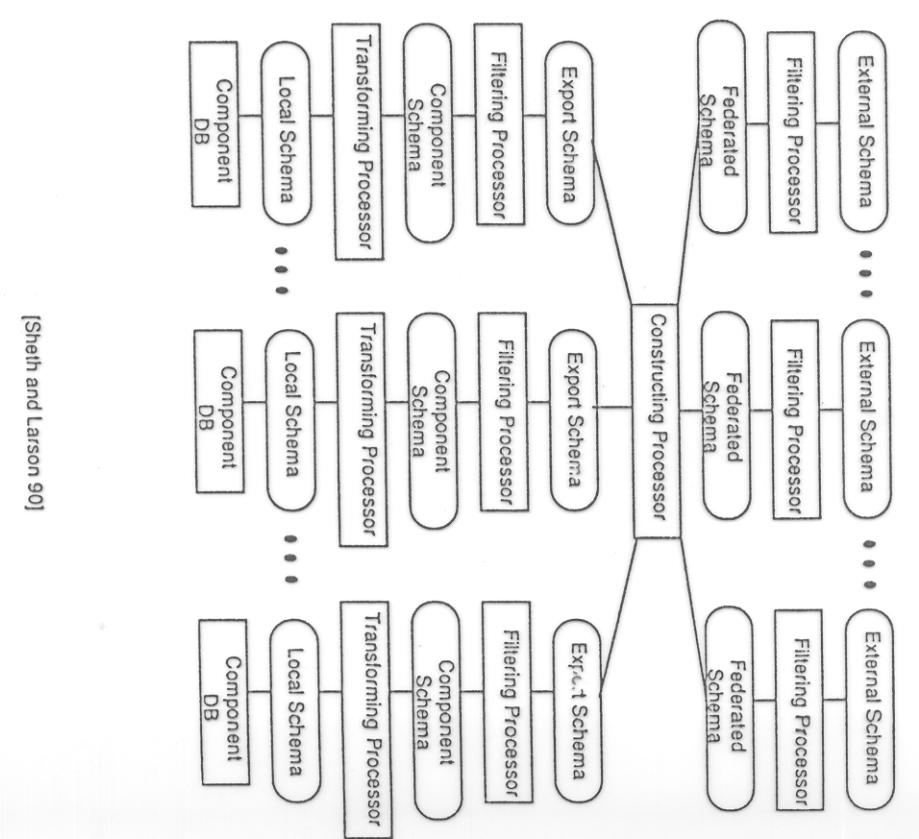
An Example FDBS Schemas  
with missing schemas at some levels



[Sheth and Larson 90]

[Sheth 88c]

## System Architecture of a FDDBS



[Sheith and Larson 90]

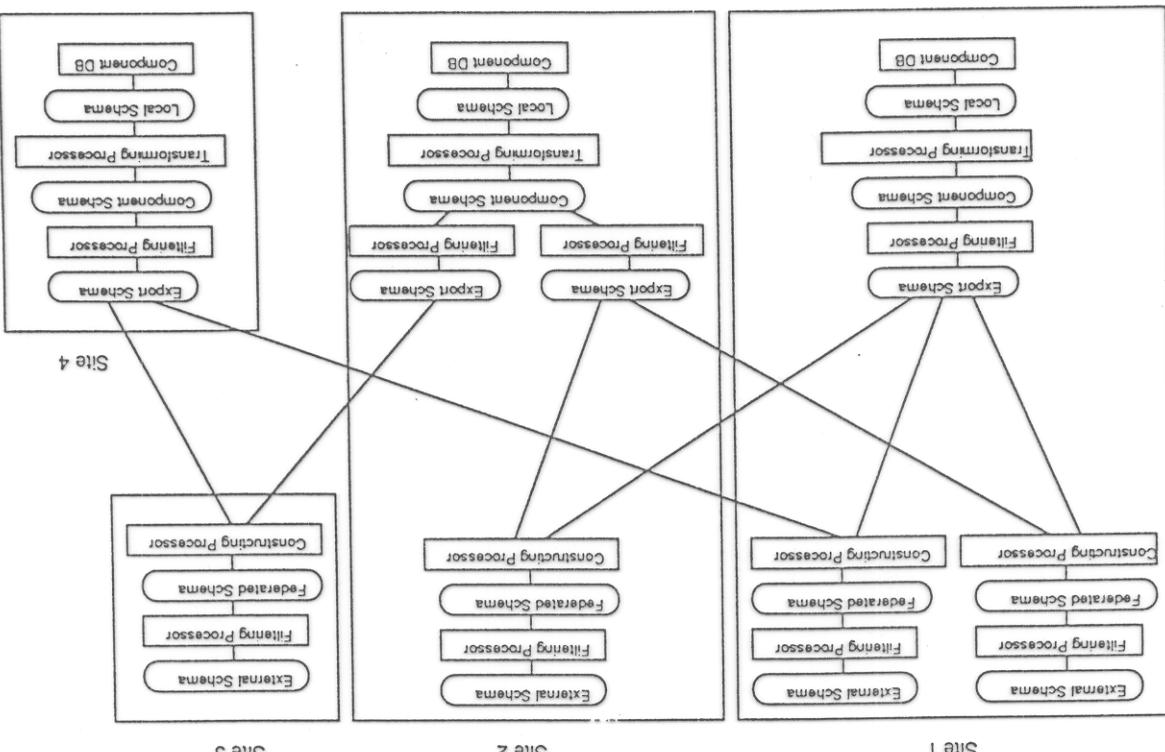
Site 4

Site 3

Site 1

## A Typical FDDBS System Configuration

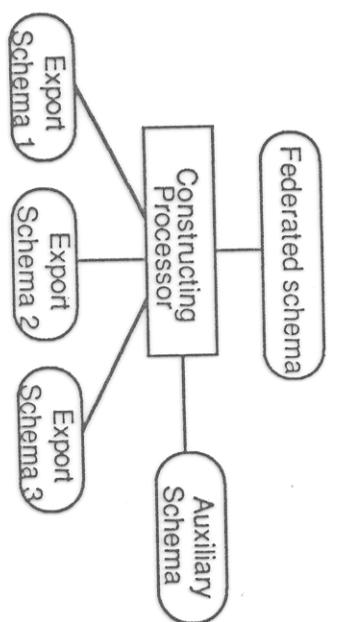
[Sheith and Larson 90]



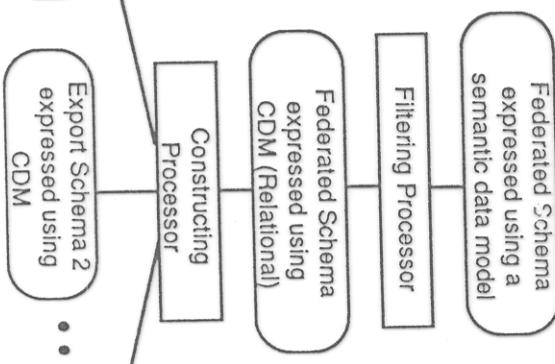
## PROCESSOR USE EXAMPLES

### FDBS B acting as a Backend to FDBS A

Using an Auxiliary Schema  
to Store Translation Information  
Needed by a Constructing Processor



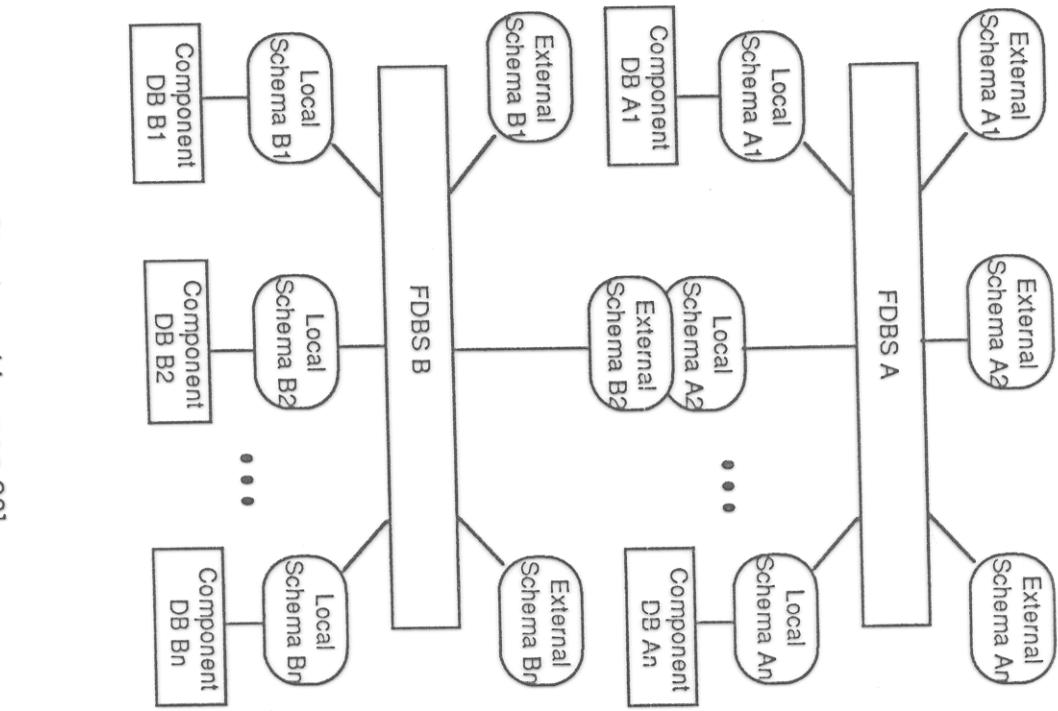
Using a Filtering Processor to Enforce Constraints  
Across Export Schemas



Export Schema 1  
expressed using  
CDM

Export Schema 2  
expressed using  
CDM

Export Schema p  
expressed using  
CDM



[Sheth and Larson 90]

## METHODOLOGY

### Three Phases

#### 1. Preintegration

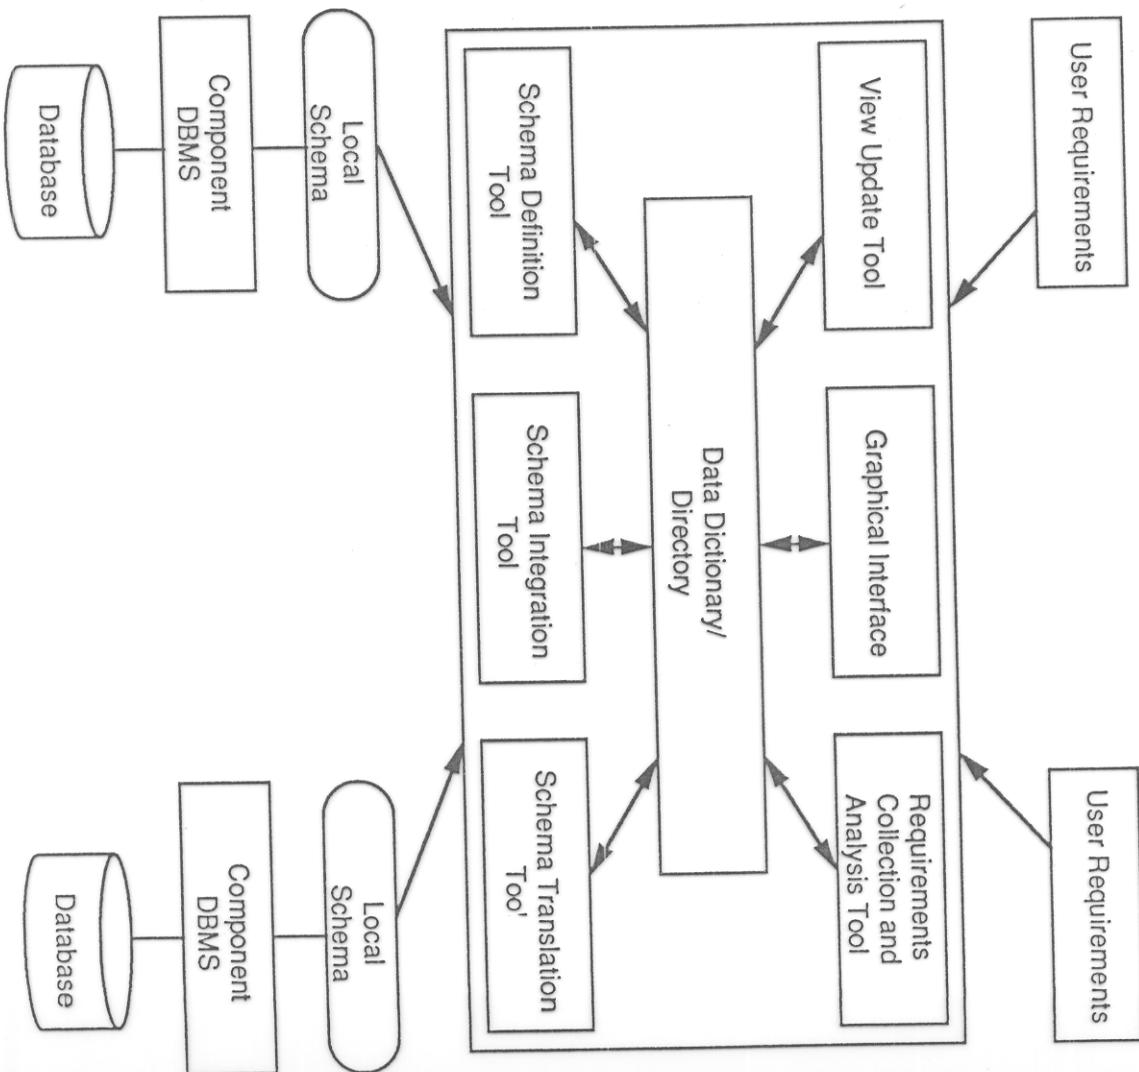
- reverse engineering

2. Design and Integration
  - schema translation
  - schema integration
  - negotiation
  - (access control)
  - (view update)
  - database design

#### 3. Operation

- multidatabase language
- operation transformation
- query processing and optimization
- transaction management
  - concurrency control and recovery
  - (access control)/security
  - (view update)

## Tools for Developing FDBSS



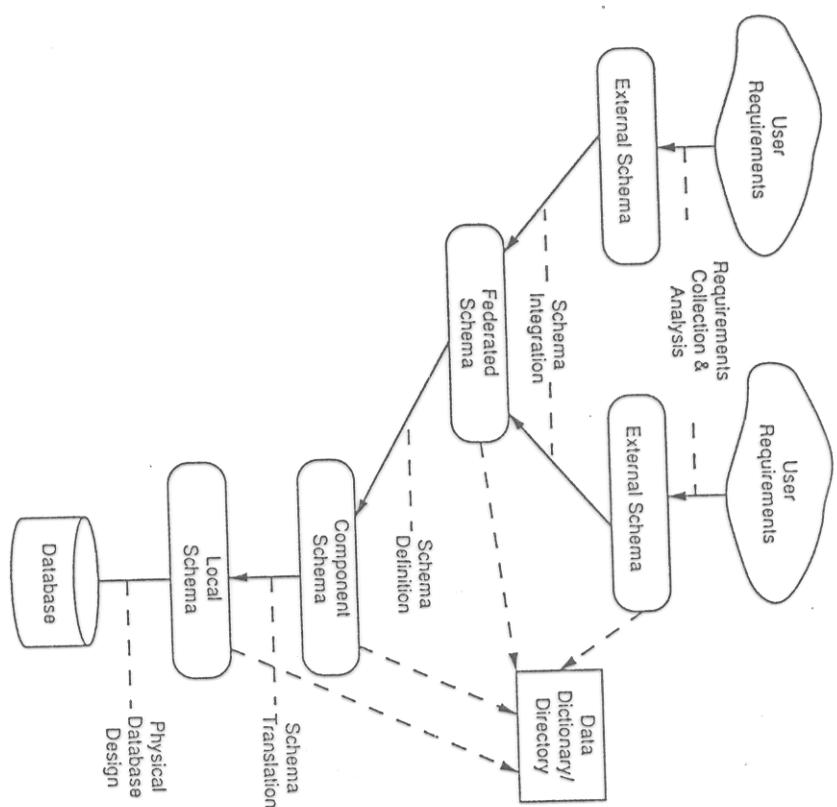
These phases need not be performed serially.  
Multiple iterations are often needed.

[Sheth 88c]

Sheth 88

## METHODOLOGY

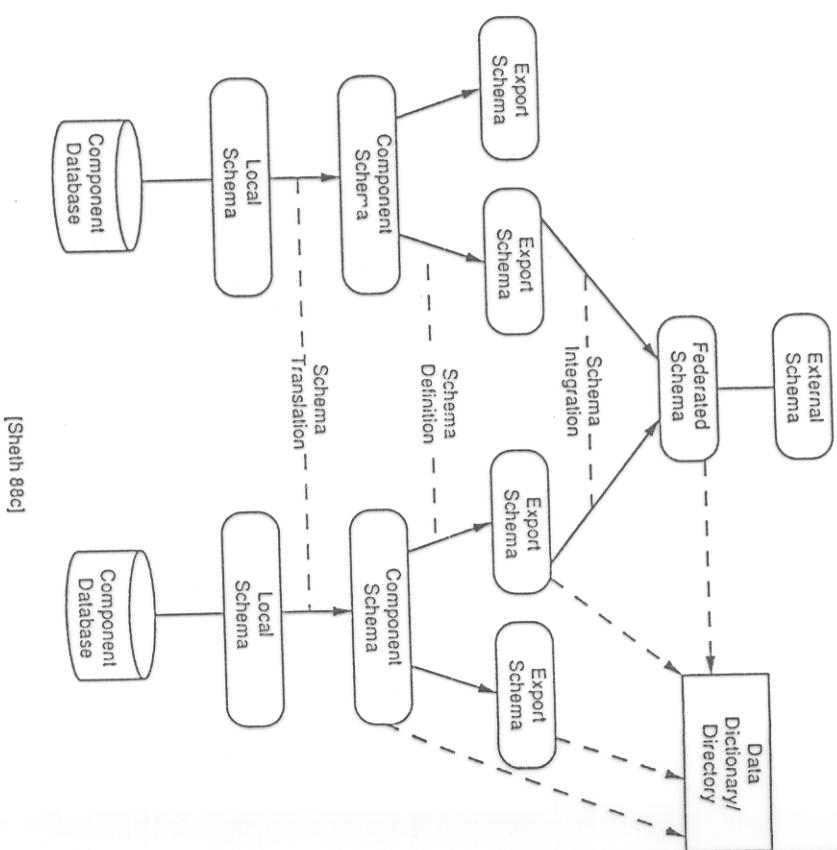
### Top-Down FDBS Building Process



[Sheth 88c]

## METHODOLOGY

### Bottom-Up FDBS Building Process



[Sheth 88c]

### 3. SCHEMA TRANSLATION

Problem: Given a source schema in one model, map it into an equivalent target schema in another model. If the source schema and target schema are according to the same data model, we have a schema restructuring mapping, otherwise we have a schema translation mapping.



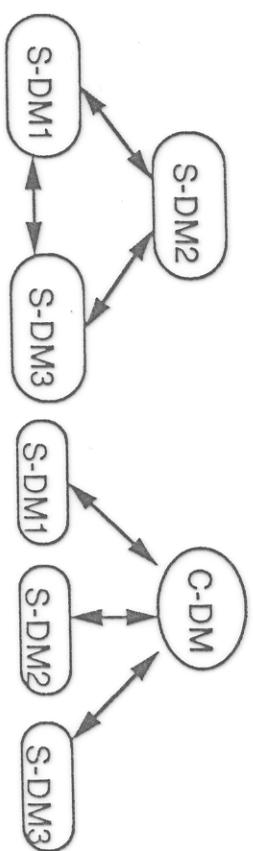
Two schemas are equivalent if

- (1) they describe the same database,
- (2) any command expressed in terms of the object of one schema can be translated into commands expressed in terms of the objects of the other schema, and
- (3) the same changes are made to the database if either source command or the target commands are executed.

Two ways to perform translation:  
Specify mapping between every pair of data models OR use a common (canonical/mapping) data model

Examples of common data models used in various schema translation efforts are:

Relational [Klug 80] [Breitbart et al 84]  
Binary [Pelagatti 78] DIAM [Senko 76]  
Extended Set Theory [Sibley and Hardgrave 77]  
E-R type [Katz 80] [Elmasri et al 85]  
Functional [Landers and Rosenberg 82]



S-DM<sub>i</sub> = Schema in Data Model "i"  
C-DM = Schema in Common Data Model

[Larson 83] [ Tsichritzis and Lochovsky 82]

## SCHEMA TRANSLATION

### Comparision Summary of Mapping:

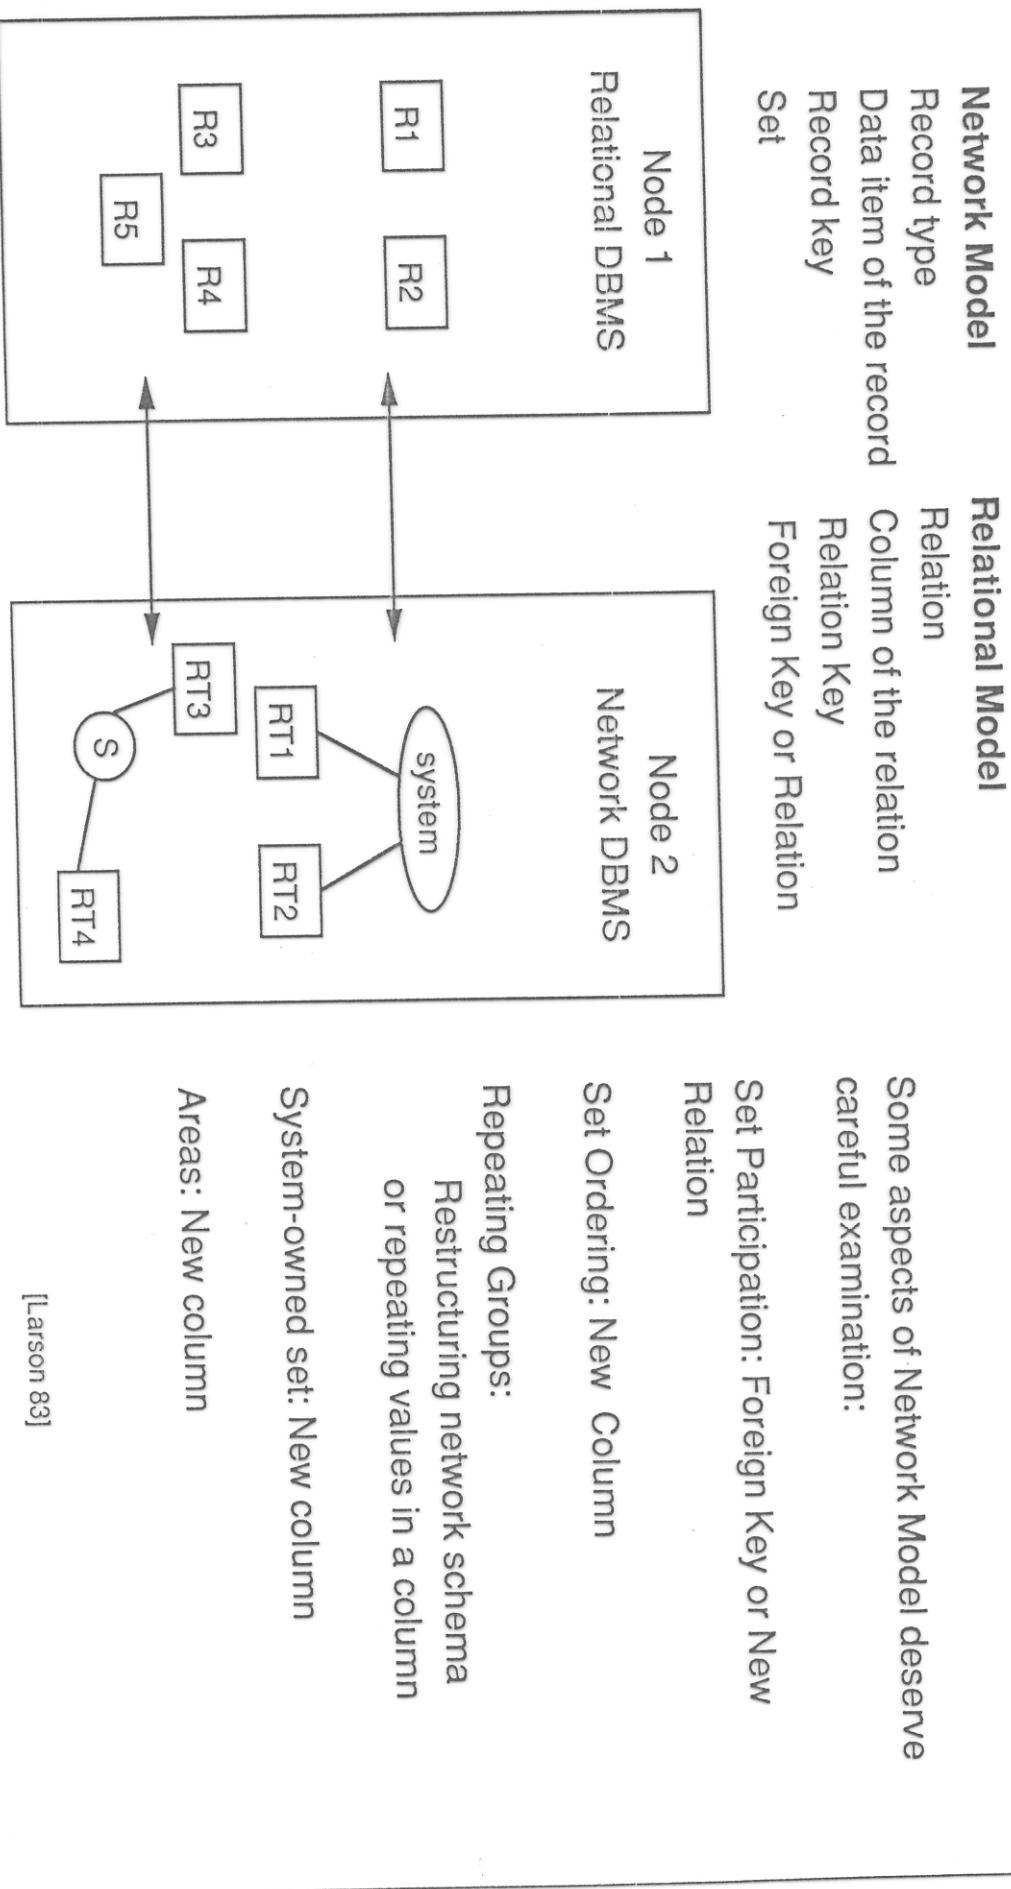
### Representation Structures of three models

	Relational	Network (DBTG)
ECR		
Entity-type	Relation	Record type
Category (lSA)	Relation	Record type and single owner single member set
Category(generalization)	Relation	Record type and single owner multiple member set
Functional relationship (1:1,!N)	Foreign key	Single owner single member set
Non-functional relationship (M:N)	Relation	Record type and owner-member sets
Single valued attribute	Attribute	Field
Multi-valued attribute	Attribute 'Nest' relation	Repeating group, vector
Entity identifier	System assigned attribute	Database key

[Elamsri et al 85]

# NETWORK TO RELATIONAL TRANSLATION

## Structural Similarities:



## Network to Relational Translation (continued)

### SCH<sub>E</sub>MA TRANSLATION

#### Remarks

(1) For each record type Ni define a relation Ri such that:

- (a) Ri contains one attribute for each data item of Ni.
- (b) If Ni has a key, the key of Ri is equal to the key of Ni; otherwise, the key of Ri is equal to the database key of Ni, which appears as an explicit attribute of Ri.

(2) For each link Lij, with owner record type Ni and member record type Nj, define relational constraints and change the existing relations such that:

- (a) The key of Ni appears as a foreign key of Rj.
- (b) One of the following sets of constraints applies, depending on the type of set membership (Ri and Rj are the relations corresponding to Ni and Nj, respectively).

- Fixed Automatic
- Mandatory Automatic
- Optional Automatic

E.g. Network to Relational  
ECR to Relational

\* More Difficult Translations:  
Non-procedural to Procedural  
Representation to Conceptual

E.g. Relational to Network  
Relational to ECR

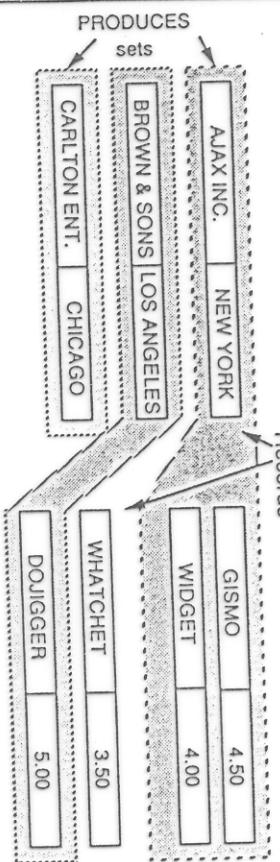
- (c) Instead of explicit constraints of (b), it is possible to derive relational schema that contains only relations and multivalued dependencies if network schema is loop-free and the links are fixed automatic.

[Tsichritzis and Lochovsky 82]  
[Vassiliou and Lochovsky 80] [Lien 80]

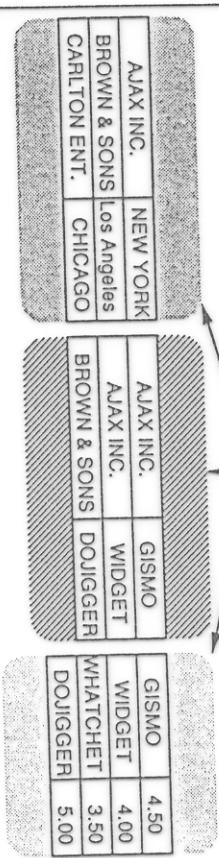
[Tsichritzis and Lochovsky 82]

## Network to Relational Translation

### An Example



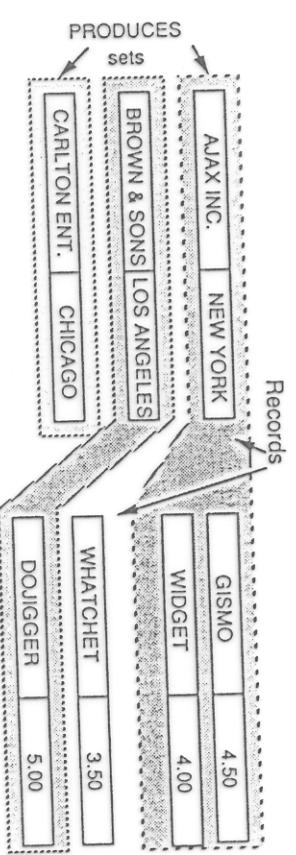
Source schema: A network database and schema.



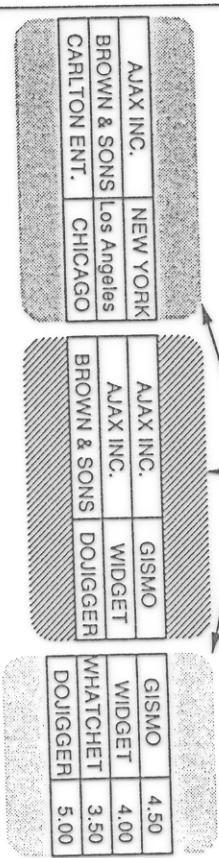
Source schema: A network database and schema.

## Network to Relational Translation

### An Example: Alternate Translation



Source schema: A network database and schema.



Source schema: A network database and schema.

Target schema: A relational database and schema.

Target schema: A relational database and schema.

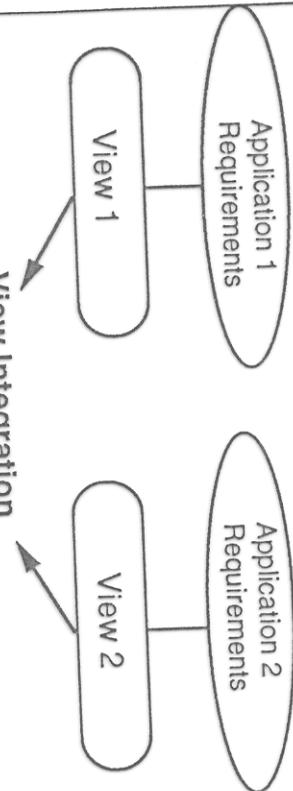
[Larson 83]

## An Example: Possible Mapping Network to Relational Translation

1. R-COMPANY.Company Name <- N-COMPANY.Company Name
  - R-COMPANY.City <- N-COMPANY.City
  - R-PRODUCT.Product Name <- N-PRODUCT.Product Name
    - R-PRODUCT.Cost <- N-PRODUCT.Cost
    - R-PRODUCT.Cost = Product-Cost
    - Key(R-PRODUCT) = Manager-Name
    - <- Key(N-PRODUCT) = Product-Name
2. Foreign-Key(R-PRODUCT.Company Name)
  - R-PRODUCT.Product Name seduced
  - Explicit Constraints = Optional Automatic
  - <- Key(R-COMPANY.Company Name)

# SCHEMA INTEGRATION

Logical Database Design:  
(*top-down design*)



Identify relationships (dependency) among objects in different schemas/databases

Develop corporate wide view of logical data

- identify redundancy and remove or maintain

Database Integration:  
(*bottom-up design*)

- increase sharing of data, provide uniform access
- maintain consistency of related data
- provide better understanding of operations

Database Integration



Database 1

Schema 1

Database 1

Schema 2

## Schema Integration Activities

## Resolving Incompatabilities

### Preintegration:

- All schemas should be in the same data model;
- Choice of Common Data Model is critical
- Specify global constraints and naming conventions

### Comparing:

- Also called Schema Analysis
- Identify schema components that represent the same information/semantics
- Define Attribute Equivalence Relationship
- Assert Object (Entity/Relationship) Equivalences

### Conforming:

- Same information should be represented using the same model constructs

### Merging:

- Components from different schemas are merged to create an integrated schema; this task can be automated

### Restructuring:

- Integrator can use operators to make changes in the integrated schema to capture semantics not used earlier in the process

### Types of Incompatibilities:

1. Differences in Models
2. Differences in Naming Conventions
3. Differences in Representations
4. Differences in Scales/Units
5. Missing Data
6. Conflicting Data Values

### Solutions

- A. Schema Translation to solve 1.
  - B. Schema Analysis and by
    - Renaming
    - Logical Restructuring
    - Scale Conversion
- These enable us to solve 2 thru 4.
- B. Auxiliary Database can be stored with auxiliary schema attached to the unified schema to solve 5 and 6.

[Landers and Rosenberg 82] [Litwin Abdellatif 86]

[Batin et al 86]

1	3	8
---	---	---

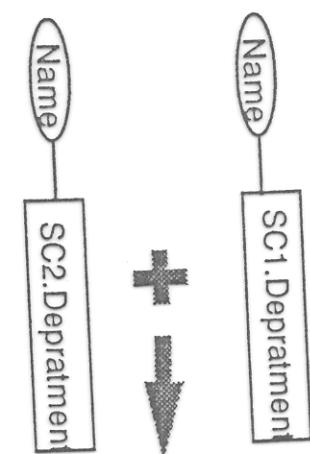
# SCHEMA INTEGRATION

## Object Relationships

### Case 3: Overlapping Domains



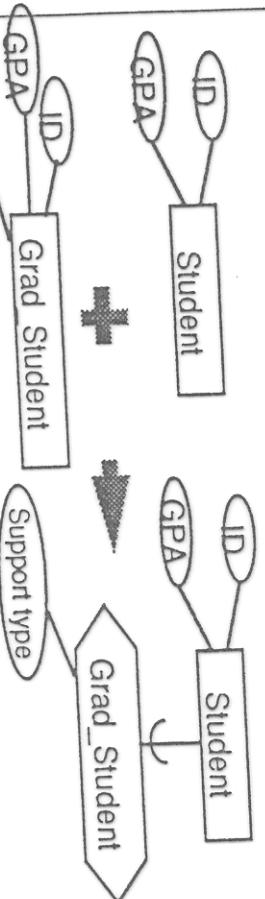
### Case 1: Identical Domains



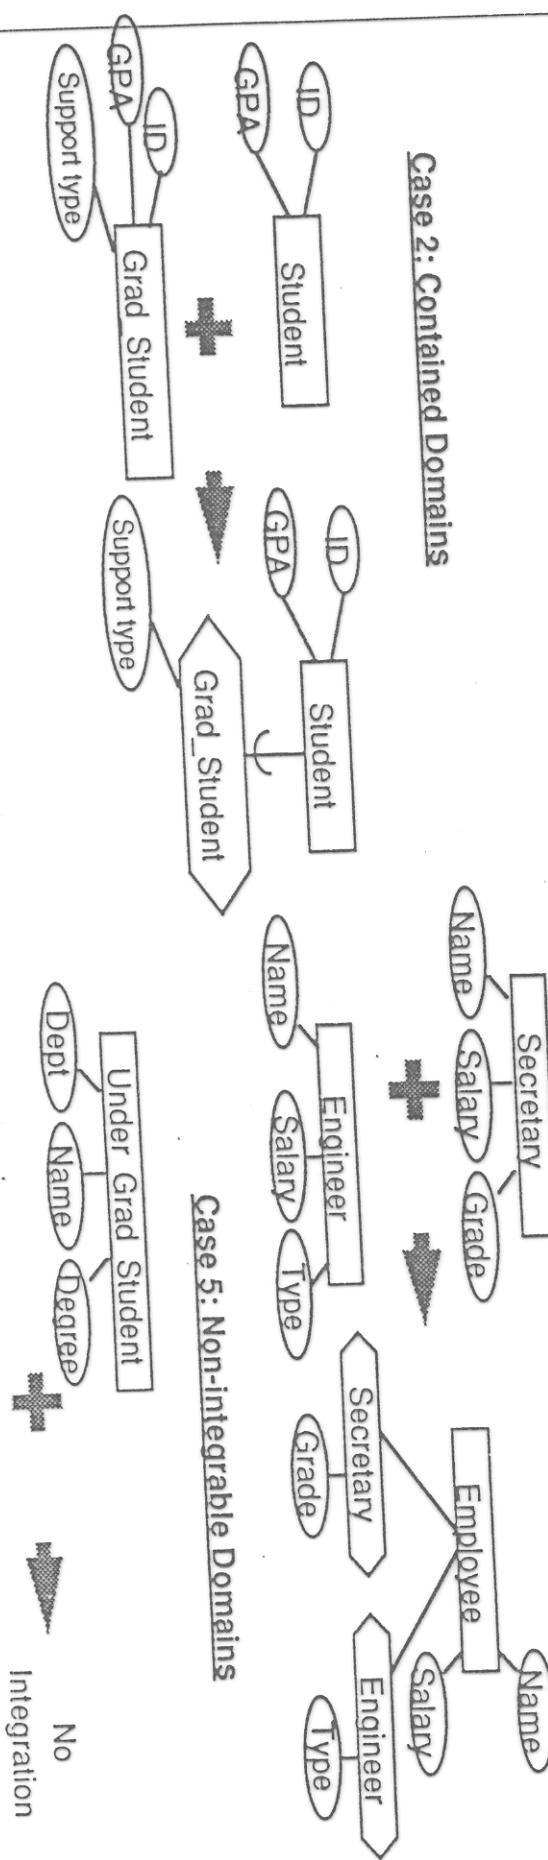
### Case 4: Disjoint Domains



### Case 2: Contained Domains



### Case 5: Non-integrable Domains



No  
Integration

[Navathe et al 86]

## WHERE IS SEMANTICS?

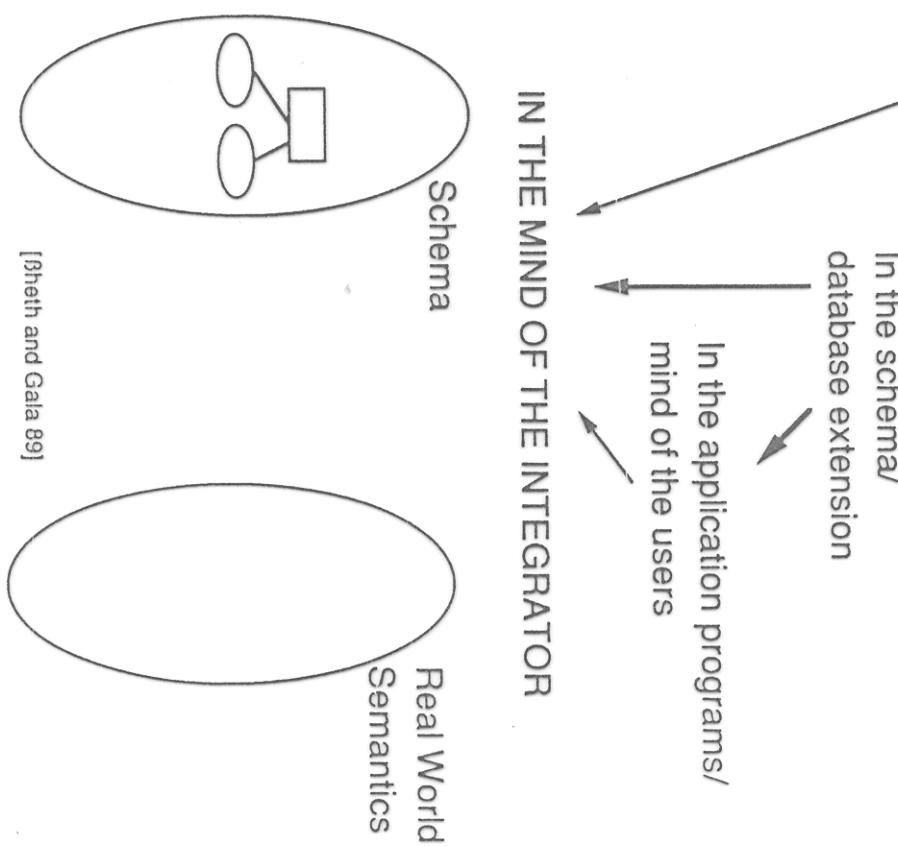
### SCHEMA INTEGRATION:

#### Observations

The key problem in integrating schemas is our **inability to completely capture the semantics** and being able to argue/reason about the semantics.

We can perform some automated reasoning about the object relationships, but **reasoning about attribute relationships cannot be automated**, it has to be done by a human.

There is no unique integrated schema for a given set of component schemas. Different semantics given by the integrator will result in different integrated schemas.



# A Case Study on Schema modeling, Analysis, and Integration

## BERDI

(Bellcore Schema Design and Integration Toolkit)

A prototype software system to develop, evaluate and demonstrate methodology and techniques to

- real world functionality
- real world scale

Some related work:

Tools: Sheth et al 88; Hayes & Ram 90; Sheth & Marcus 92  
Techniques & Concepts: Battini et al 86; Battini et al 91; Collet et al 91; Dayal & Hwang 84;  
Elmasri, et al 86; Navathe & Gadgil 82; Spaccapietra et al 91, ...

# Solving the right problem: A perspective on Schema Integration

Different perspectives, different needs:

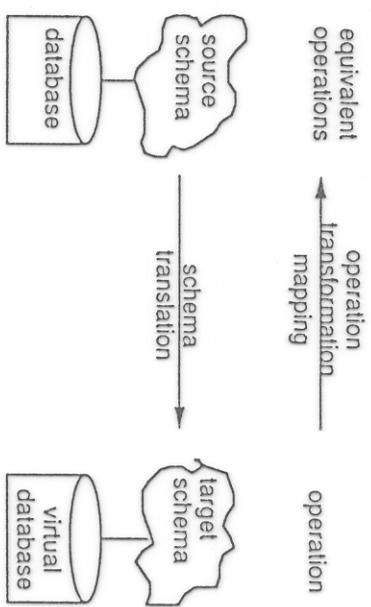
- Bottom-up (schema) integration-- To develop integrated schema(s) to give transparency to the distributed/heterogeneous databases?
- Top-down (view) integration --To integrate views/subschemas of the same application to develop the database schema for that application?
- To analyze schemas for dependencies among them to facilitate interoperability?
- To develop (the schema of) a new application that encompasses one or more existing applications?
- To perform a **shallow** analysis/integration (e.g., without attributes, functions, or behavior), and/or **deep** analysis/integration resulting in complete dependency/integration mapping?

## OPERATION TRANSFORMATION

Query Translation  
Relational to Network

Given a source schema, a database according to the source schema and an equivalent target schema, map operations on target schema to equivalent operations on source schema. If the source and target schemas are according to the same data model then we have view mapping, otherwise operation transformation mapping. Operation transformation

mapping is used to translate a query operation on a target schema into equivalent operations on the source schema.



Transforming nonprocedural relational data manipulation command into a sequence of procedural network commands:

Select =>  
A looping statement in network data manipulation language in which each record is retrieved and examined. Records satisfying the selection condition are moved to the target relation.

```

T = Select (R1,condition) => r = get_first(N1);
Loop: if condition(r);insert r
      r = get_next(N1); go to
            
```

Loop

Project =>  
A looping statement in which each record is retrieved and the specified items are copied to the target table.

Join =>  
Two nested loops in which the outer loop reads successively records from one record type and the inner loop successively reads records from second record type in the set owned by each of the record of the first record type.

## OPERATION TRANSFORMATION

View Mapping  
(through Query Modification)

### DATABASE RELATIONS:

<i>EMPLOYEE</i>		<i>DEPARTMENT</i>	
NAME	ADDRESS	DEPT	LOCATION
Smith	15 Bloor	Paint	Paint King
Jones	25 King	Toy	Toy Queen
Carter	171 Yonge	Paint	Pipe King

### VIEW DEFINITION:

```
EMPVIEW <-
SELECT
  EMPLOYEE.NAME, ADDRESS, MANAGER
  FROM EMPLOYEE, DEPARTMENT
  WHERE EMPLOYEE.DEPT = DEPARTMENT.DEPT
```

### VIEW:

```
EMPVIEW
NAME      ADDRESS   MANAGER
Smith    15 Bloor   Miles
Jones   25 King    Quin
Carter  171 Yonge  Miles
```

### Now a query on EMPVIEW:

ANSWER <-

SELECT MANAGER

FROM EMPVIEW

can be rewritten into the following query

by a technique called Query Modification:

ANSWER <-

SELECT MANAGER

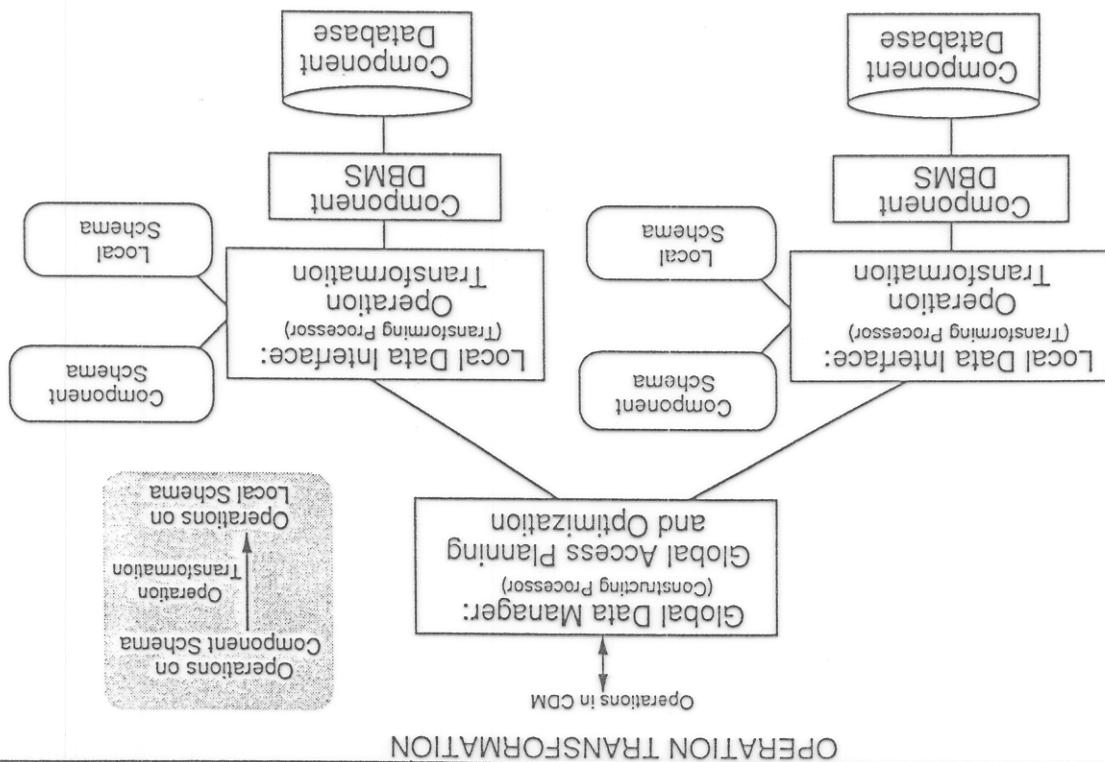
FROM DEPARTMENT

WHERE DEPARTMENT.DEPT = SELECT

EMPLOYEE.DEPT

```
FROM EMPLOYEE
WHERE NAME = 'Smith'
```

[Stonebraker 75] [Tsichritzis and Lochovsky 82]



## OPERATION TRANSFORMATION

### Query Translation

#### Relational to Network

#### Remarks

Transforming nonprocedural relational data manipulation command into a sequence of procedural network commands:

Select =>

A looping statement in network data manipulation language in which each record is retrieved and examined. Records satisfying the selection condition are moved to the target relation.

Operation Mapping can be established by direct mappings or by using intermediate mapping data model as in schema translation.

Following complexities exist in query translation:

```
T = Select (R1,condition) => r = get_first (N1);
      Loop: if condition(r) insert r into T
            r = get_next(N1); go to Loop
```

Project =>  
A looping statement in which each record is retrieved and the specified items are copied to the target table.

representation  
(This may be an important reason for not choosing a procedural model as a model.)

Join =>

Two nested loops in which the outer loop reads successively records from one record type and the inner loop successively reads records from second record type in the set owned by each of the record of the first record type.

[Larson 83] [Vassiliou and Lochovsky 80]

## 6. QUERY PROCESSING AND OPTIMIZATION

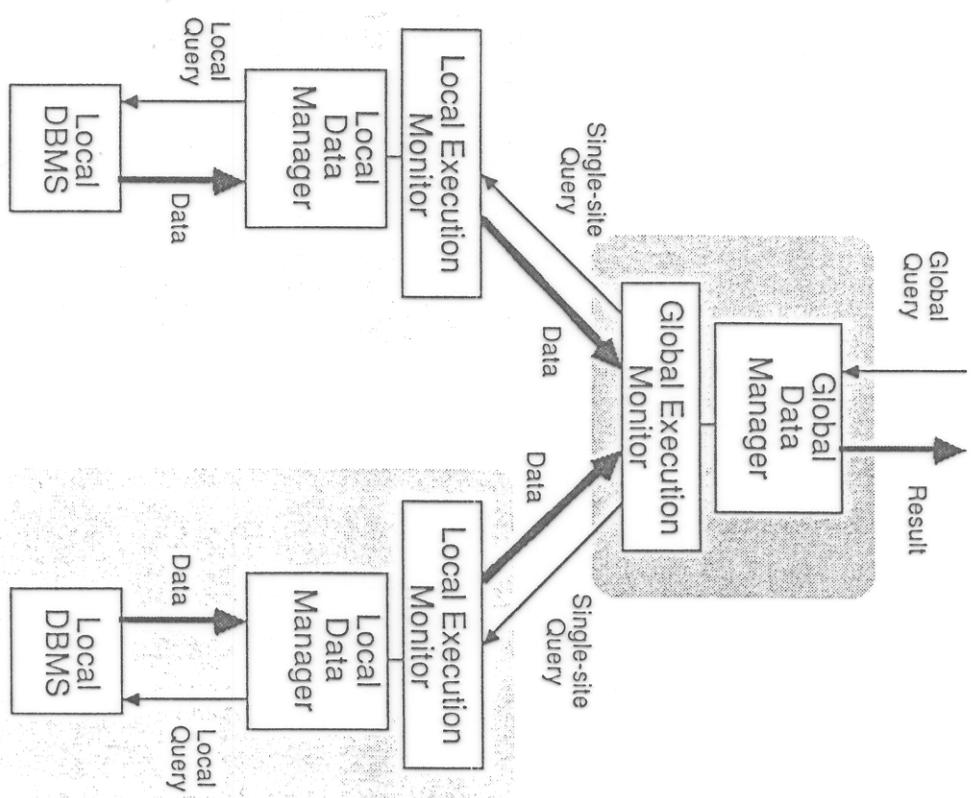
### Query Processing

**Goal:** Utilizing the potential for parallel processing in a distributed system and minimizing processing cost by reducing the amount of data moved between sites.

**Issues:** Additional complexities introduced by the heterogeneous environment.

- (1) Cost of performing a local query may differ greatly from site to site.
- (2) The ability to receive and reference "moved" data may vary from site to site. Many DBMS do not support creation and loading of temporary databases.
- (3) Local querying capabilities may be limited.
- (4) Local DBMSs vary in their capabilities to optimize.

[Landers and Rosenberg 82]



## Global Optimization

## Global Optimization

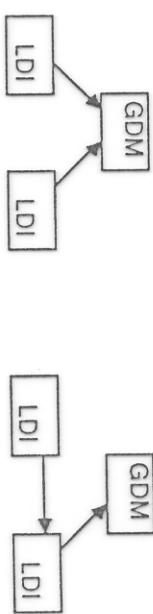
There is a trade-off between the amount of work done by the GDM and the software complexity of LDI. There is also a trade-off between amount of communication and processing done by different sites. Depending on how a global transaction is divided, a range of alternatives is possible.

Examples of alternatives are:

- (1) More work by GDM, more communication but simple:  
Translate the global query in smallest possible single site subqueries. Each site may be sent several subqueries. Collect partial result at GDB and merge.
- (2) Medium work by GDM, medium communication, medium complexity of GDM and LDI's:  
Translate the global transaction into largest possible single site subtransactions (one transaction for each site where data is located). This will reduce communication and needs less work when merging partial results. However, the LDI's have to be sophisticated.

(3) Less work by GDM, less communication, complex GDM and LDI's:  
Generate efficient programs such that LDI's participate in global optimization. Partial results can be sent to GDM or other LDI's. LDI's have to be able to support some additional functionalities like sorting, removing duplicates, handling and merging temporary files.

E.g., consider multisite join:



## Local Optimization

## Local Optimization

Goal: Optimize processing of the subtransaction received from GDM against the local DBMS.

Opportunities for Optimization:

- Local Query Language
- Physical Database Organization
- Access Path Selection

Sequences of joins, selects, and projects can be optimized.

E.g., a select followed by a project can be combined into a single loop that selects records satisfying the select condition and copies the projected items.

[Onuegbe et al 83] [ Dayal and Goodman 82]

# TRANSACTION MANAGEMENT AND CONCURRENCY CONTROL

## CONSISTENCY CRITERIA

A transaction is a set of operations on database. A traditional transaction has following ACID properties:

1. **ATOMICITY:** Either all or none of the operations should "happen". It commits or aborts.
2. **CONSISTENCY:** The transaction should be a correct state transformation. This is typically left to the programmer.
3. **ISOLATION:** Each transaction should see a correct picture of the database state, even if concurrent transactions are executing.
4. **DURABILITY:** Once a transaction commits, all its effects must be preserved, even if there is a failure.

### Local and Global Consistency: Conflict Serializability

Serializability is the well known criteria used to define "correct" executions of concurrent transaction.

Current research seeks to address the following question: Which of these properties are appropriate or required in a FDBS and how can they be selectively relaxed? What are useful transaction models? How do autonomy and heterogeneity affect transaction management?

**Weak Serializability:** final state serializability, view serializability, quasi serializability

**Mutual Consistency:**  
Single Copy Serializability  
Eventual Serializability

[Haeder and Reuter 83]

## CONCURRENCY CONTROL

There are three methods of Concurrency Control:

- \* Locking
- \* Timestamping
- \* Optimistic

**Locking:** If a transaction accesses a data item, it sets a lock on it. Another transaction with a conflicting lock request has to wait until the lock is released.

**Timestamp:** Each transaction get a timestamp at the begining and the conflicting transactions are processed according to the timestamp.

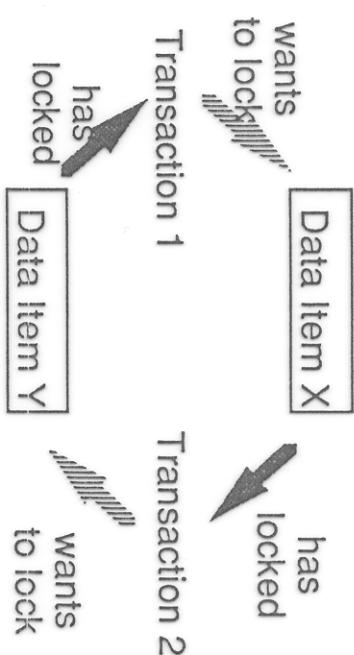
**Optimistic:** After a transaction is executed, it checks if it could have conflicted with other transaction. If so it "aborts" otherwise it "commits".

Two transactions conflict if one of them write a data item while other reads or writes the same data item.

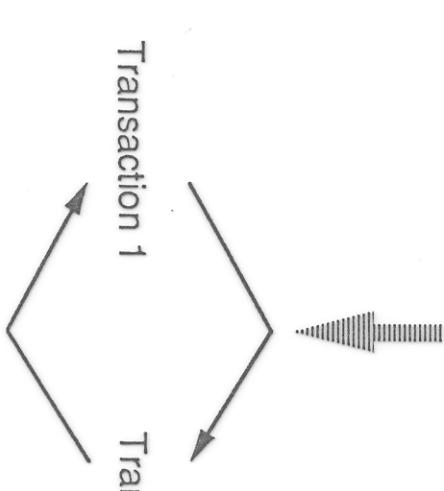
[The above definitions are incomplete and oversimplified.]

[Bernstein and Goodman 81]

Locking is the most popular method.  
But, whenever a locking is used,  
deadlock  
can occur.



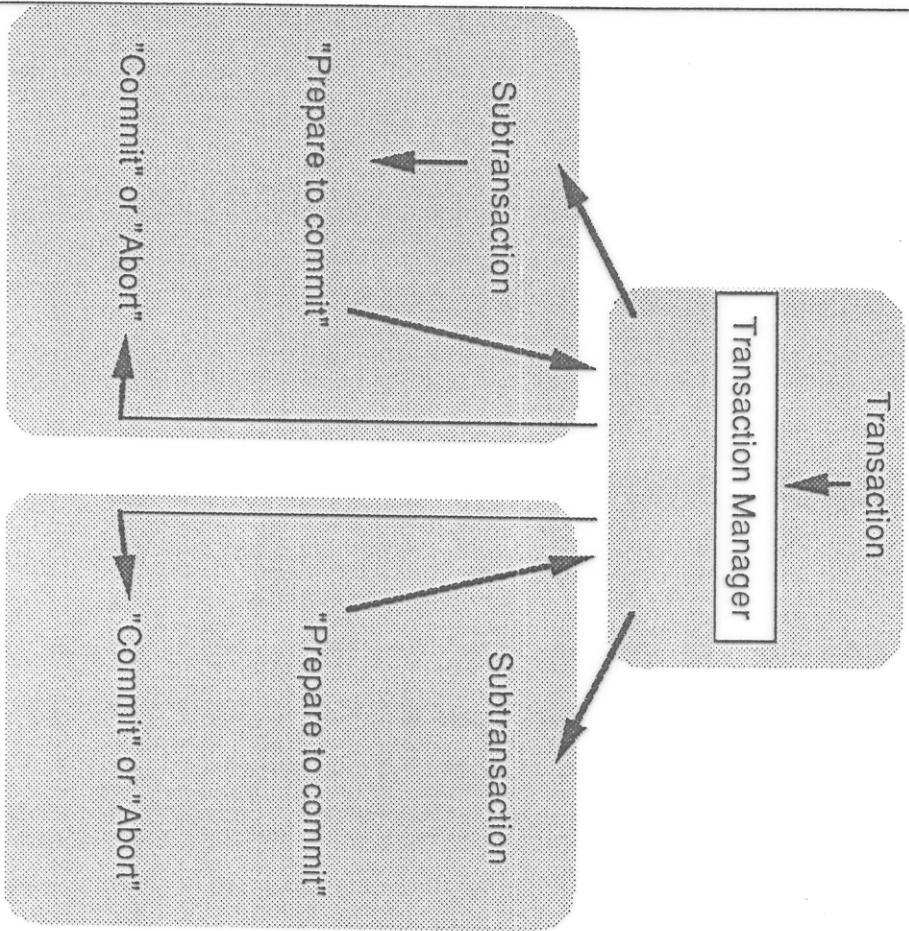
Arrows show wait-for graph. A cycle in a wait-for graph implies a deadlock.



## TRANSACTION MANAGEMENT

### EFFECT OF HETEROGENEITY

In order to preserve atomicity property of the transaction, TWO-PHASE COMMIT protocol is used.



1. Different DBMSs use different Concurrency Control methods. E.g., if the global TM uses a two phase locking mechanism and a component DBMS uses a timestamp ordering, the latter cannot provide the lock point information required by the former.

✓2. Even if the Concurrency Control methods are the same, implementations may vary significantly.

E.g., locking methods may vary in deadlock handling methods, number of items locked by one lock (locking granularity) may vary, etc.

3. The DBMSs may vary in commit protocols.  
E.g., Some may not provide two phase commit, while some may.

[Bernstein and Goodman 81]

[Logar and Sheth 86]

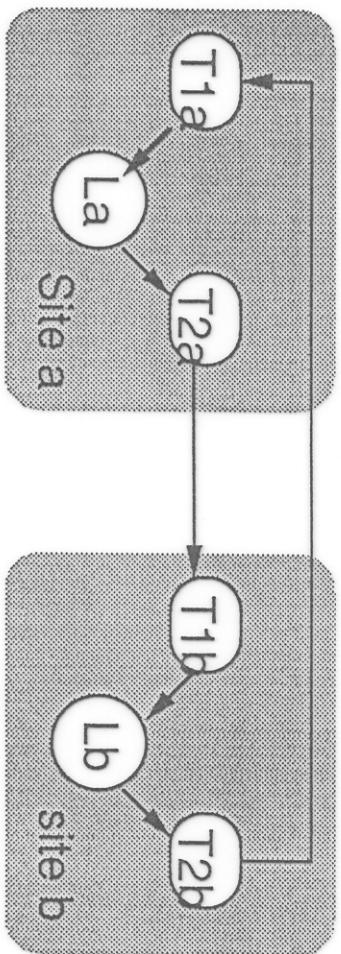
## EFFECT OF AUTONOMY

### Global Deadlock Detection Problem

A global deadlock involves transactions at more than one site. Problem of global deadlock detection is very difficult because (1) a local process (part of local DBMS or Operating System) does not know about non-local transactions, and (2) a global process does not know about the local transactions.

Most practical solution seems to be time-out.

Example of global deadlock:



1. Design Autonomy leads to all problems listed in "Effect of Heterogeneity".
2. Even if a component DBMS has the information required for global TM, communication autonomy may prevent the global TM from accessing it.
3. If the global TM has detected a conflict, execution autonomy of a component DBMS may prevent it from taking actions (e.g., block or restart a transaction) necessary to resolve the conflict.

[Du et al 89b]

[Logar and Sheth 86]

## EFFECT OF AUTONOMY

### Direct Conflict vs Indirect Conflict

## The Transaction Processing Architecture for HDBMS

Consider Component Databases/DBMSs: D<sub>1</sub>: a; D<sub>2</sub>: b,c

Global Transactions: G<sub>1</sub>: W<sub>G1</sub>(a) R<sub>G1</sub>(c)  
G<sub>2</sub>: R<sub>G2</sub>(a) W<sub>G2</sub>(b)

Local Transaction at D<sub>2</sub>: L: R<sub>L</sub>(b) W<sub>L</sub>(c)

Local Execution E<sub>1</sub> at D<sub>1</sub>: W<sub>E1</sub>(a) R<sub>E2</sub>(a)

Local Execution E<sub>2</sub> at D<sub>2</sub>: W<sub>E2</sub>(b) R<sub>E1</sub>(b) W<sub>L</sub>(c) R<sub>G1</sub>(c)

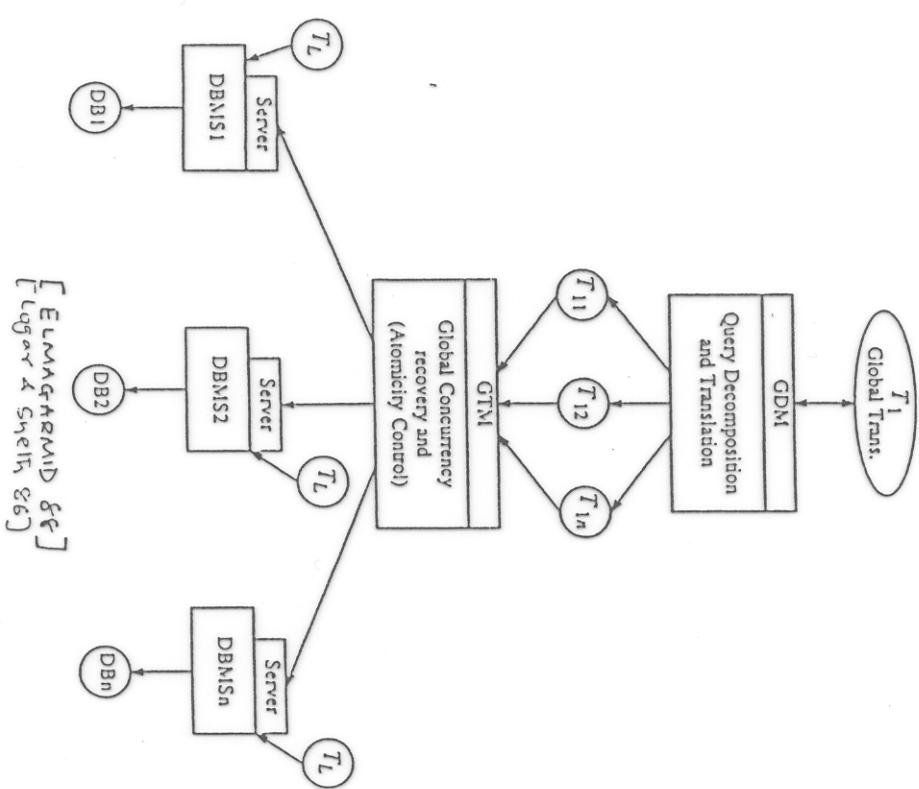
In E<sub>1</sub>, G<sub>1</sub> and G<sub>2</sub> directly conflict,  
in E<sub>2</sub>, G<sub>1</sub> and G<sub>2</sub> indirectly conflict through the  
operations of L.

### Execution Order vs Serialization Order

The serialization order between two operations might be different than their execution order. The indirect serialization order of R<sub>G1</sub>(c) and W<sub>G2</sub>(b) in E<sub>2</sub> remains unchanged even if the execution order is changed to E<sub>3</sub>.

E<sub>2</sub>:: W<sub>G2</sub>(b) R<sub>L</sub>(b) W<sub>L</sub>(c) R<sub>G1</sub>(c)  
E<sub>3</sub>:: W<sub>L</sub>(c) R<sub>G1</sub>(c) W<sub>G2</sub>(b) R<sub>L</sub>(b)

[Du et al 89]



## TRANSACTION MANAGEMENT AND CONCURRENCY CONTROL

### Alternatives

1. Allow only read (retrieval) operation. Do not allow write (update) operation. Concurrency control problem does not arise if all transactions are read only.
2. Assume all DBMSs provide locking and support two-phase commit. Detect deadlocks by time out.
  - DBSSs do not report the local serialization order of transactions.
3. Limit concurrency. (e.g., allow one global transaction per component DBMS; let the locking granularity be single site and check acyclicity of the site graph [Logar & Sheh 86][Alonso et al 87]; limit local transaction to be read only or write only on non-replicated data. [Breitbart and Silberschatz 88]
4. Compromise autonomy, or require info for global CC, or enforce specific serialization order. [Sugihara 87], [Pu 87]? Also DDTs, top-down approach [Elmagarmid & Du 90] to enforce specific serialization order..
5. Different consistency criteria or relax strict serializability. E.g., Quasi serializability [Du and Elmagarmid 89] (allows only serial execution of global subtransactions). S-Transactions [Eliassen & Veijalainen 87], Sagas [Garcia-Molina & Salem 87]. Value date [Litwin & Tirri 88]. Flexible transaction [Elmagarmid et al 90].
6. Prevent different schedules of conflicting transactions at different sites without loss of autonomy. Ticketing method [Georgakopoulos et al 91].

## Problems in Enforcing Global Serializability

- Transaction execution and serialization orders can be different.
- DBSSs do not report the local serialization order of transactions.
- To determine the global serializability, the MDBS must take into account conflicts caused by local transactions.
- Information about local transactions is not available.

## Problems in Enforcing Global Serializability

Example: Serial Global Execution & Strict Local Schedules

## Problems in Enforcing Global Serializability

Example: Serial Global Execution



site A:

$r_{G_1}(a)$   $w_{G_2}(a)$

$G_1 \rightarrow G_2$

site B:

$w_{G_1}(b)$   $r_{G_2}(c)$

$r_{T_{local}}(b)$

$w_{T_{local}}(c)$

$G_2 \rightarrow T_{local} \rightarrow G_1$

Georgakopoulos

## Analogous Execution and Serialization Orders

### Rigorous Schedules and Schedulers

**Definition:** Let  $S$  be a serializable schedule. We say that transactions in  $S$  have *analogous execution and serialization orders* if for any pair of transactions  $T_i$  and  $T_j$  such that  $T_i$  is committed before  $T_j$  in  $S$ ,  $T_i$  is also serialized before  $T_j$  in  $S$ .

**Definition A** *transaction management mechanism is rigorous if the following two conditions hold: (i) it guarantees strictness and (ii) no data item may be written until the transactions that previously read it commit or abort.*

**FACT 1:** Strict schedules do not assure that transactions have analogous execution and serialization orders.

**FACT 2:** Schedulers allowing all schedules in which transactions have analogous execution and serialization orders are expensive, and difficult to design and build.

**Theorem:** *If a transaction management mechanism ensures rigorously, it produces (conflict) serializable schedules in which transaction execution and serialization orders are analogous.*

[BGRS91]  
Georgakopoulos

Georgakopoulos

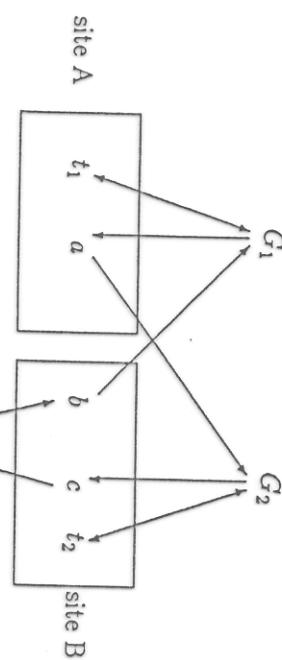
## Ticket Methods

Determining the local serialization order in non-rigorous systems

- The subtransactions of global transactions take a ticket from the LDBS they access.

The current value of the ticket is stored in the LDBS.

- Tickets are logical timestamps. The actions to read ( $r(ticket)$ ) and increment the ticket value ( $w(ticket+1)$ ) are part of each subtransaction.



**Theorem:** *The tickets obtained by the subtransactions of global transactions determine their relative serialization order.*

- Tickets create direct conflicts between the subtransactions at each LDBS.
- Subtransactions are serialized according to the ticket values they obtain.
- If not, they get aborted by the local concurrency control.

$$\text{siteA : } r_{G_1}(t_1) w_{G_1}(t_1 + 1) w_{G_1}(a) r_{G_2}(t_1) w_{G_2}(t_1 + 1) r_{G_2}(a), \text{ i.e.,}$$

$$G_1 \rightarrow G_2$$

$$\text{siteB : } w_{T_{local}}(b) r_{G_1}(t_2) w_{G_1}(t_2 + 1) r_{G_1}(b) r_{G_2}(t_2) w_{G_2}(t_2 + 1)$$

$$w_{G_2}(c) r_{T_{local}}(c), \text{ i.e.,}$$

$$G_2 \rightarrow T_1 \rightarrow G_1$$

[GRS91]  
Georgakopoulos

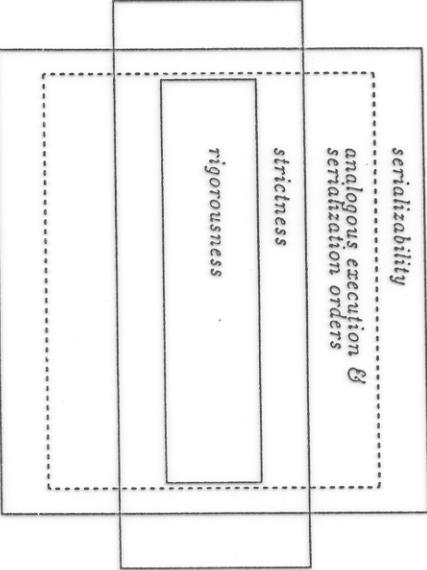
[GRS91]  
Georgakopoulos

## The Effects of the Take-A-Ticket Approach

Example

## Rigorous Schedules and Schedulers

### Relation of Rigorousness and strictness



## Rigorous Transaction Management

### Example: The The Implicit Ticket Method (ITM)

How ITM determines the local serialization order of each subtransaction?

The commitment order (implicit ticket) of the subtransactions at each DBS determines their relative serialization order.

How ITM enforces global serializability?

For any pair of global transactions  $G_i$  and  $G_j$ , it is sufficient to commit the subtransactions of  $G_i$  before the subtransactions of  $G_j$  or vice versa.

Theorem: *ITM guarantees global serializability if the following conditions are satisfied:*

- Transaction management.
  - Hierachical composition of autonomous schedulers.
- Class of rigorous schedulers includes:  
strict 2PL, conservative TO, and OCC using serial validation

[BGRS91]  
Georgakopoulos

## Using Tickets to Enforce Global Serializability

Example: The Optimistic Ticket Method (OTM)

OTM performs SGT for every global transaction  $G$ :

1. Creates a node for  $G$ .
2. For every recently committed transaction  $G_c$  adds an edge:

$G \rightarrow G_c$       if at any site  $\text{ticket}(G) < \text{ticket}(G_c)$ , or

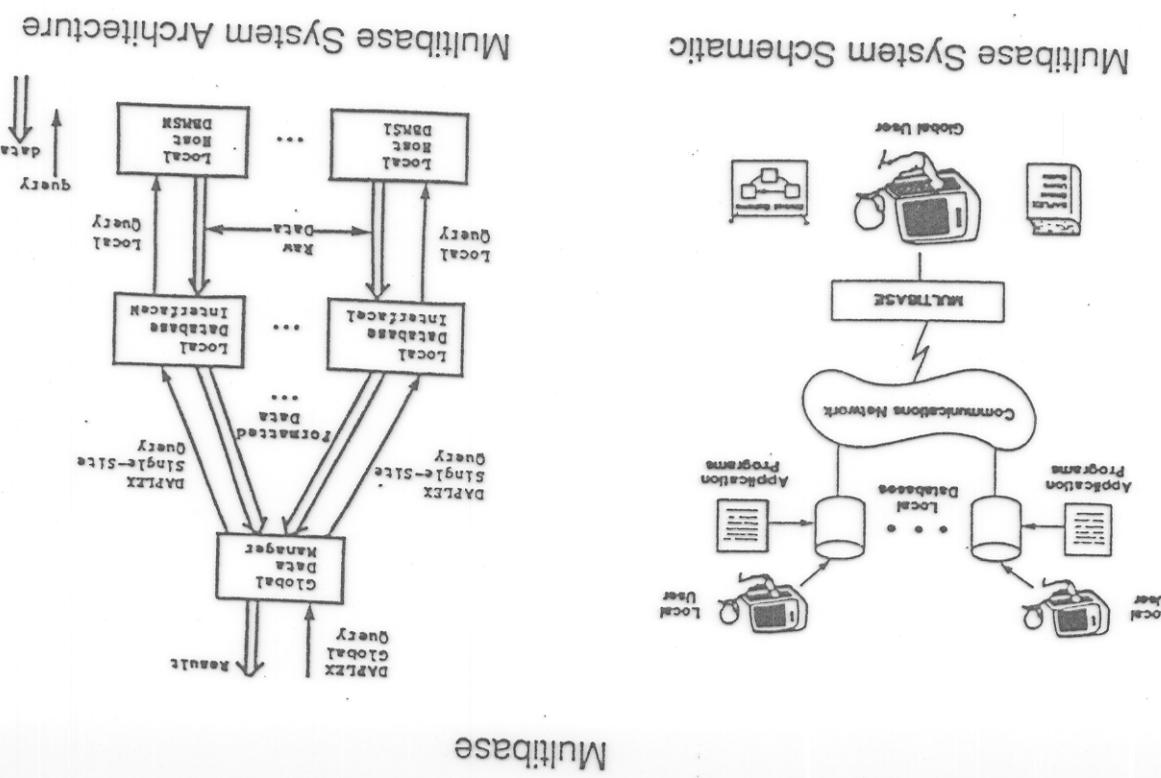
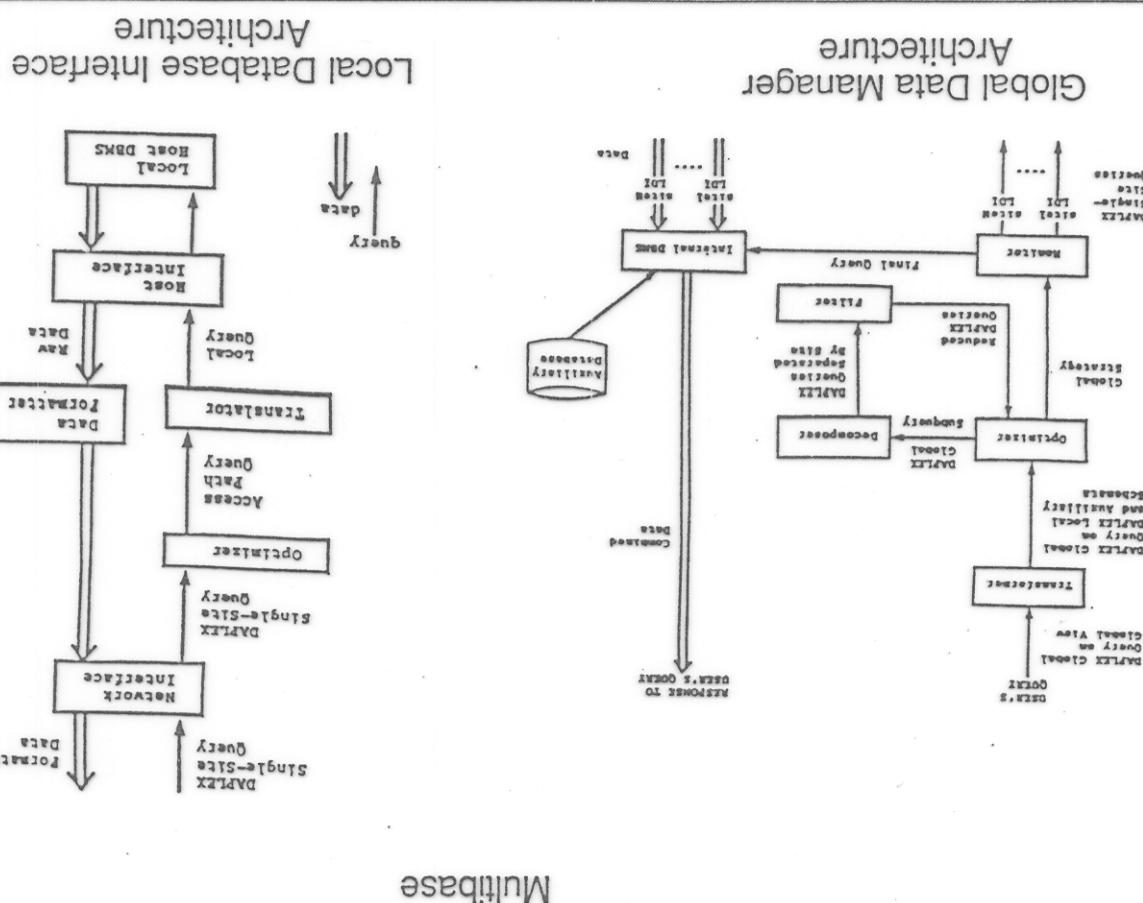
$G \leftarrow G_c$       if at any site  $\text{ticket}(G) > \text{ticket}(G_c)$

3. If there is no cycle,  $G$  is validated and is committed.

Theorem: OTM guarantees global serializability if the following conditions are satisfied:

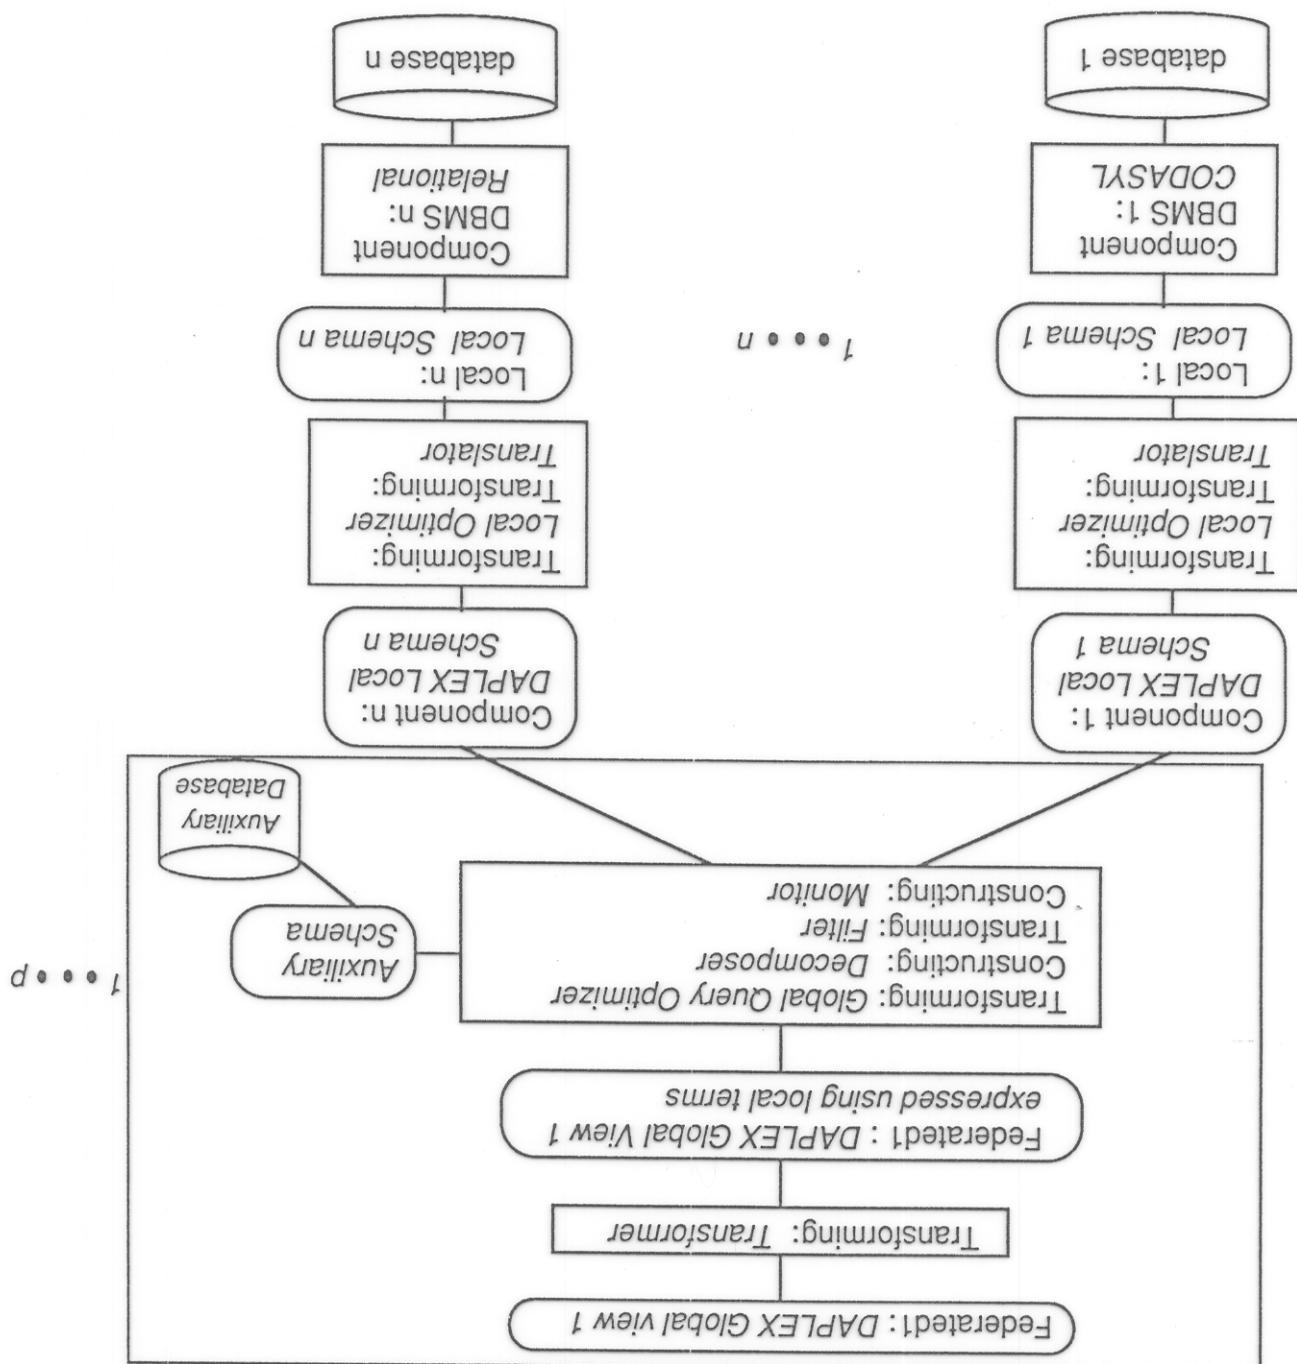
1. The LDBS assure local serializability.
2. Each global transaction has at most one subtransaction at each LDBS.
3. Each subtransaction has a visible prepare to commit state.

[GRS91]  
Georgakopoulos



[Sheth and Larson 90]

[Lundes and Rosenthal 82]



## Multibase Architecture

## DDTS Important Features

Started in 1979 at Honeywell.

Geographically Distributed (Minneapolis, Phoenix)

Schema Architecture:

Heterogeneous Hardware, Operating System,  
DBMSS.

Five Level Schema Architecture  
(see section 2.4, page 31)

System Architecture:

Application Processor  
Data Processor  
(see pages 28 and 29)

Communication: DSA  
(MAP in a second similar system)

Common Data Model: Relational

Conceptual Data Model: E-C-R

Current System Configuration:

Four Nodes:  
Three Honeywell DPS-6 with IDS-II DBMS  
One DPS-8 with RAM DBMS  
Also ties to another HDDBMS called  
Shop Floor Two Level Query System  
(Three Node, DDTS like system with  
different Hardware, OS, DBMSS)

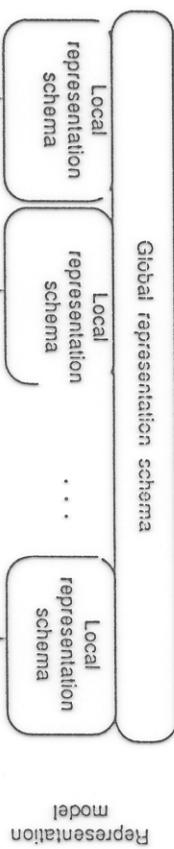
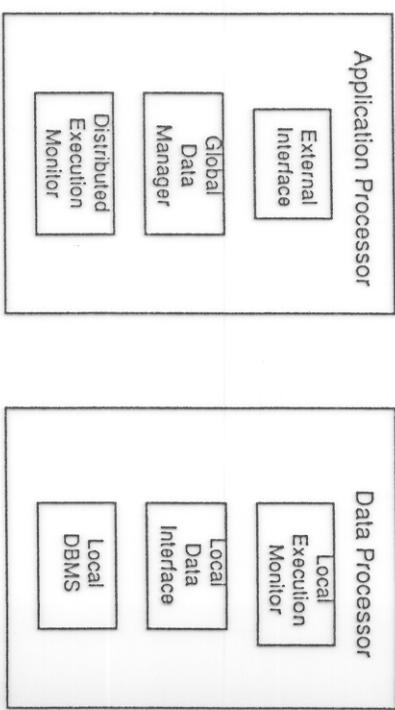
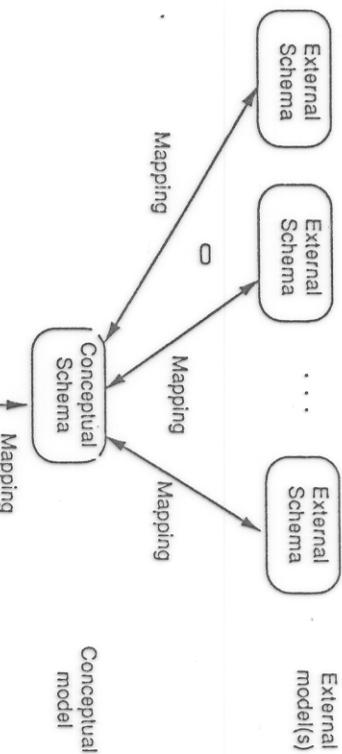
Compiled Queries can be stored.  
Allowed updates in past. Changed Operating

System for Concurrency Control.

Lacks good performance because of  
prototype nature.

[Dwyer and Larson 87]

## System Architecture of DDTs: Schemas in DDTs Terminology



External Interface (E) translates a set of commands (transaction) on the conceptual schema to a transaction on the canonical (representation) schema. It may also enforce semantic integrity constraints.

Global Data Manager (GDM) translates a transaction into subtransactions for different sites at which required data reside. This function is also called global access planning. It may perform global optimization. GDM is also called Materialization and Access Planner.

Distributed Execution Monitor (DEM) coordinates executions of subtransactions at different sites.

Local Execution Monitors (LEM) manages subtransaction evaluation at the local site and interacts with DEM.

Local Data Interface(LDI) translates the subtransaction in canonical model into a set of commands in data model or local DBMS.

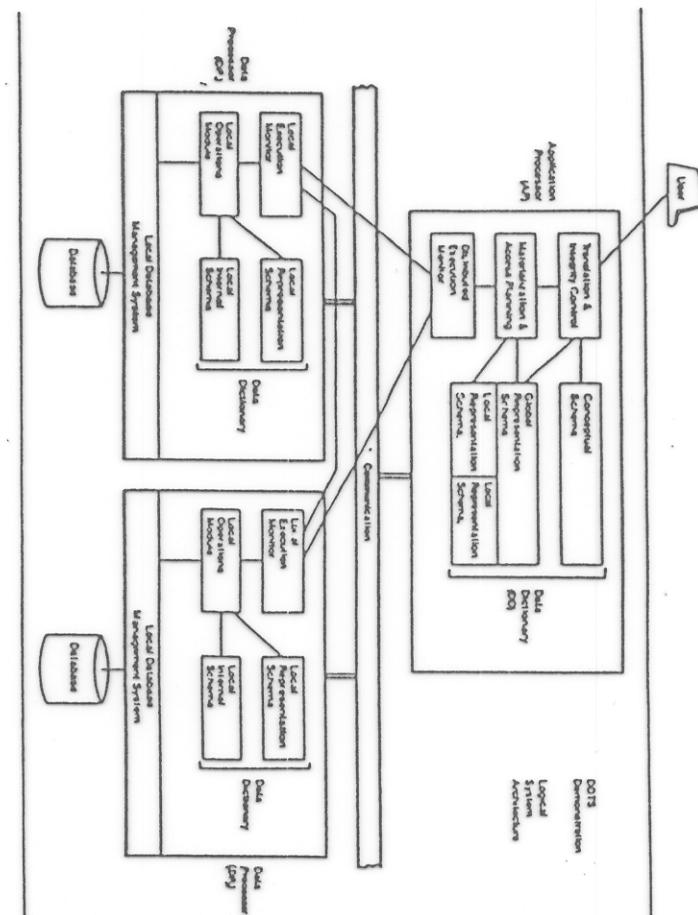
Adding local and global representation schemas allow better representation of the data distribution and DBMS heterogeneity, as well as provide data independence and multiple views. Many systems use same model at representation and conceptual level, thus degenerating to a four level schema architecture.

[Devor et al 82]

## System Architecture of DDTs: Processors in DDTs Terminology

## System and Schema Architecture

DDTS

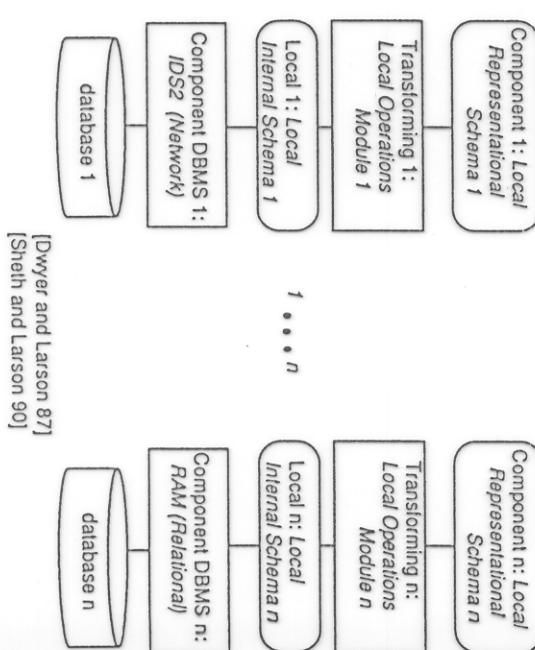
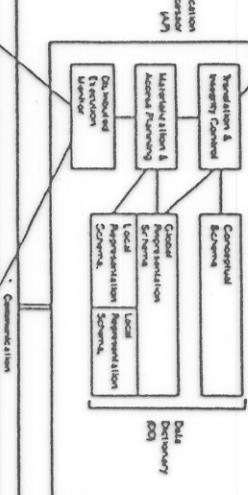


## DDTS Architecture

External: Conceptual Schema

Transforming: Translation Filtering: Integrity Control

DDTS  
Decomposition  
Logical  
Schema  
Architecture



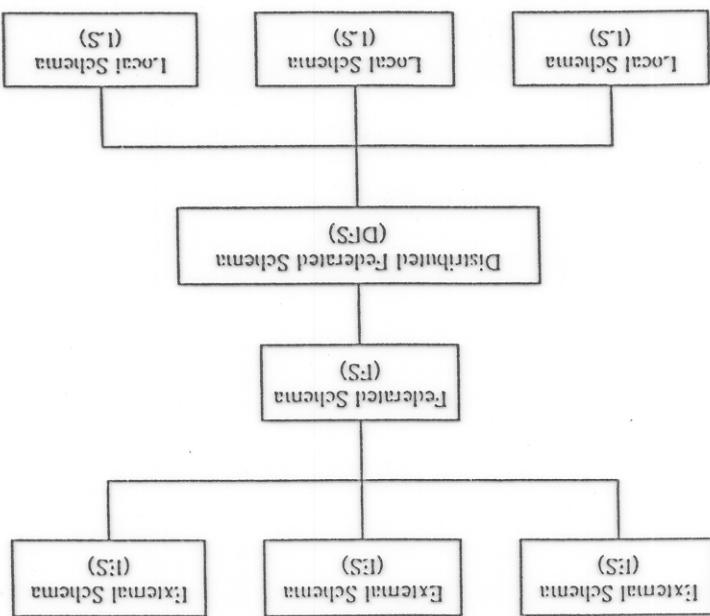
## MERMALD: Important Features

- Provides the illusion of a single DBMS
- Accesses existing databases through the existing DBMS
- Provides an integrated view of the multiple underlying databases
- A user may express a query in a standard language using the relations in the integrated view
- MERMALD determines how to answer the query
- MERMALD sends subqueries to each DBMS site
- Translators translate the subquery to the DBMS query language
- The answer is integrated into a common report
- Support for multiple federated schemas
- Extensive query optimization

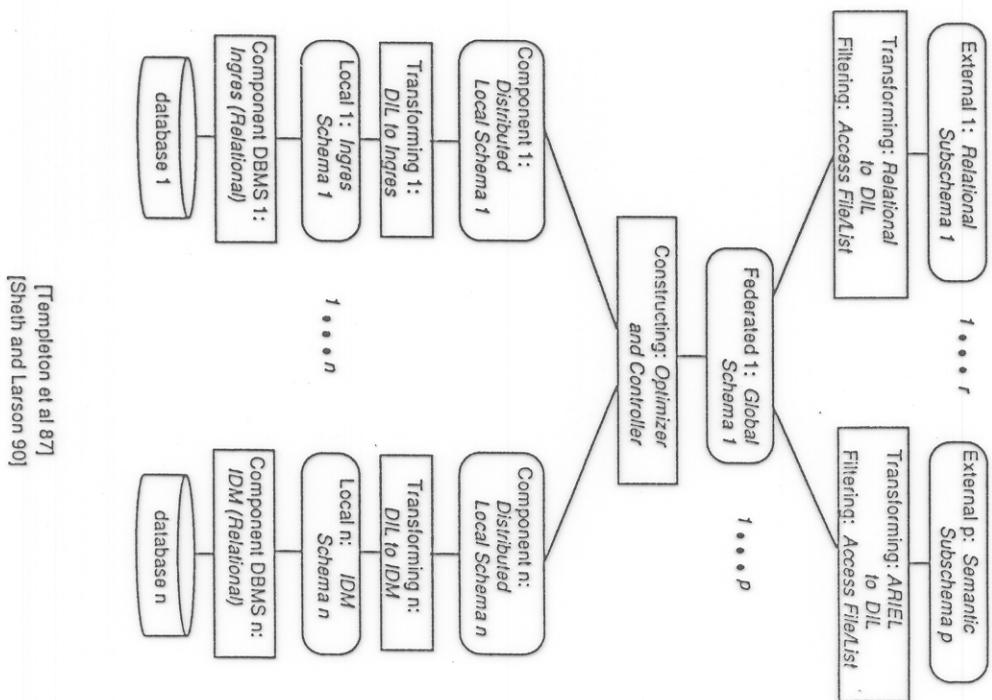
## CASE STUDIES

### MERMALD

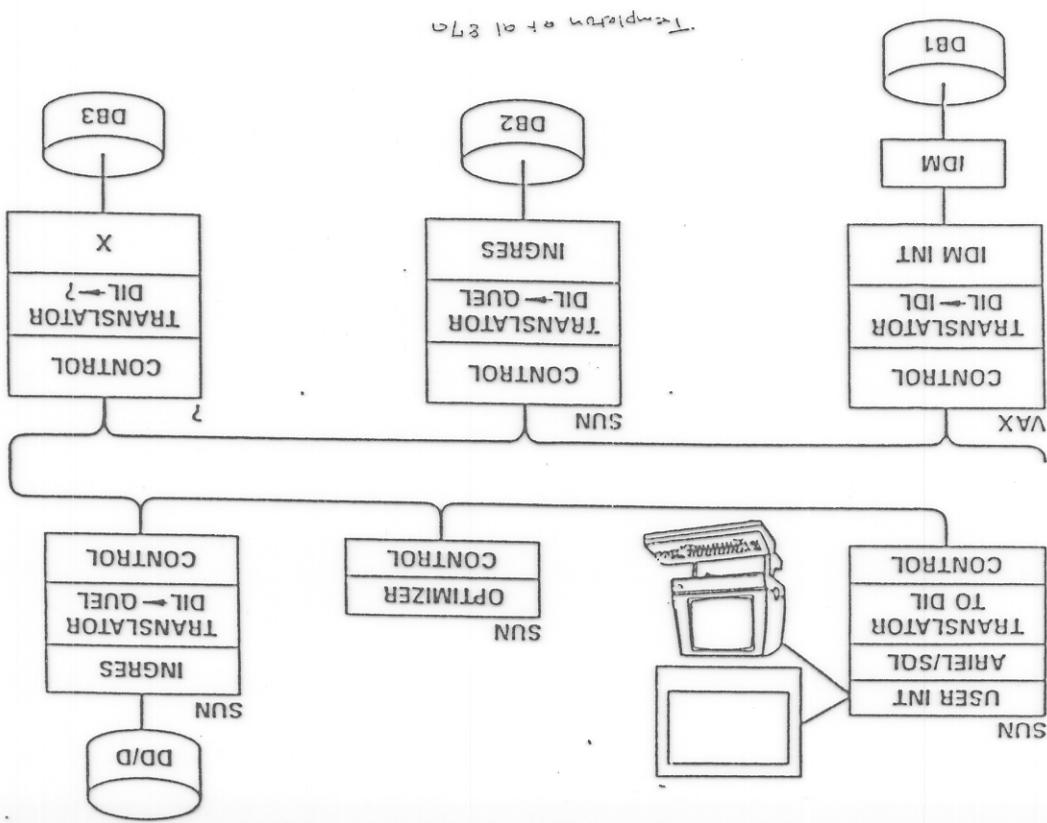
#### SCHEMA ARCHITECTURE



## Mermald Architecture



[Templeton et al 87]  
[Sheth and Larson 90]



MIRMAID: SYSTEM ARCHITECTURE

## MRDMS Important Features

Datamodel homogeneous, but semantically heterogeneous

No global schema but (a) system provides functions for manipulating data belonging to distinct/nonintegrated schemas, and (b) users face a variety of views of reality

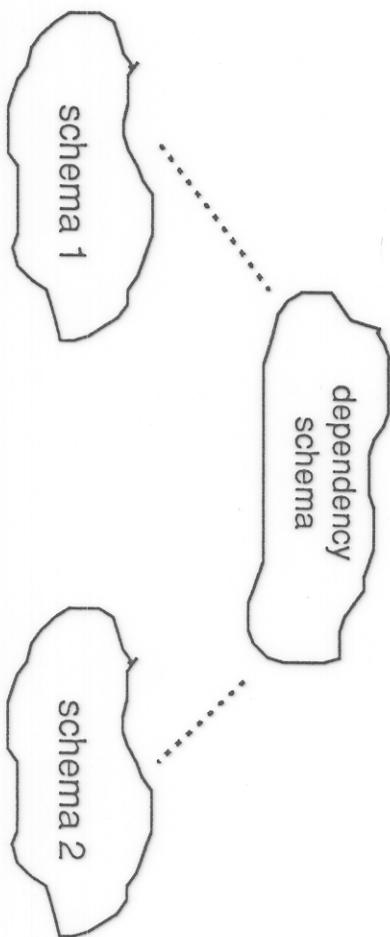
Functionality:

A. DBA can define dependency schema which describe interdatabase relationships in terms of privacy, integrity, data meaning, etc.

- joining of data in different database schemas
- user defined value types for dynamic transformation of attributes
- interdatabase queries
- dynamic aggregation of data, etc.

Interesting features:

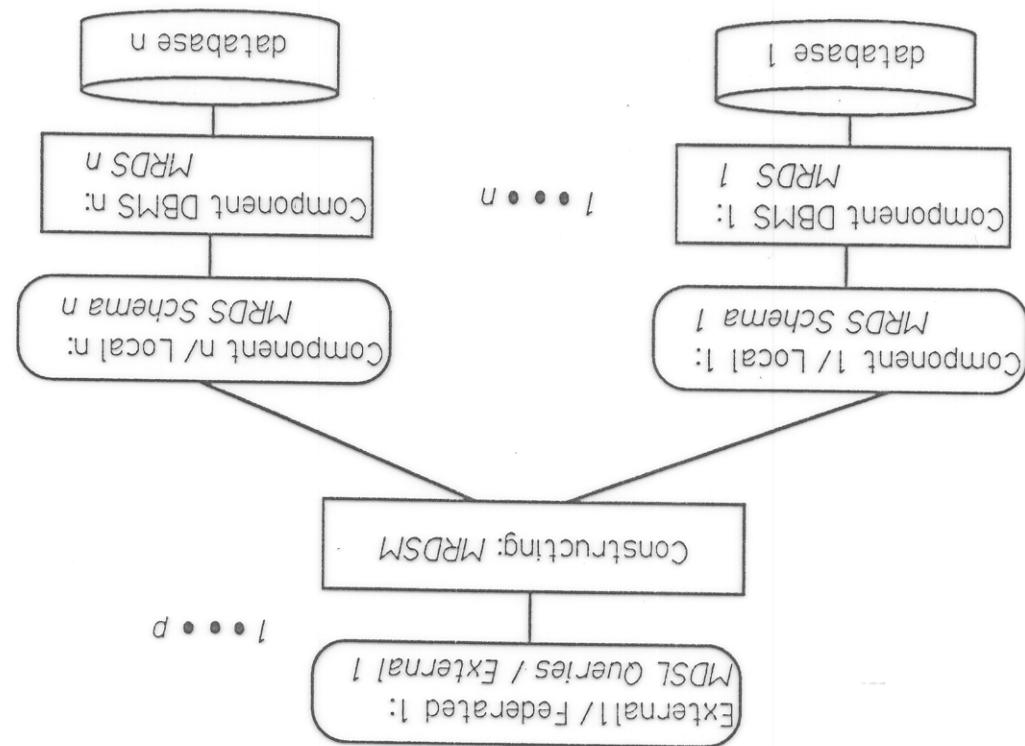
- multiple identifiers
- semantic variables
- dynamic attributes
- interdatabase queries



[Litwin and Abdellatif 86]

[Schem of Lesson 9]

### MRDSM Architecture



Sheet A Lasson 40

	Effort	Types of Components	CDM	DBMS	Reference
	Significant Features				
ADDs	Tightly coupled, limited updates, in-house use of the prototype system	More like loosely coupled, intermediate language, interactive, menu-driven user interface, data dictionary editor/browser, in-house use of the prototype system	More like loosely coupled, intermediate language, interactive, menu-driven user interface, use of local and long haul communication optimization, integrity constraint checking, local query processing system	Tightly coupled architecture, integrity constraints must be compiled, forms-based user interface like tightly coupled, queries must be compiled, forms-based user environment, attention to autonomy	IISs
CALIDA	Tightly coupled updates, in-house use of the prototype system	More like loosely coupled, intermediate language, interactive, menu-driven user interface, data dictionary editor/browser, in-house use of the prototype system	Tightly coupled, compiler, global query optimizer, auxiliary schema, IBM control, global optimization, extensions to include texts, experience with communication systems, extensions to include texts, access to implementation, use of local and long haul communication optimization, integrity constraints must be compiled, forms-based user interface like tightly coupled, queries must be compiled, forms-based user environment, attention to autonomy	Mermaid	
DDTs	Tightly coupled architecture, integrity constraints must be compiled, forms-based user interface	More like loosely coupled, integrity constraints must be compiled, forms-based user interface like tightly coupled, queries must be compiled, forms-based user environment, attention to autonomy	Tightly coupled, compiler, global query optimizer, auxiliary schema, IBM control, global optimization, extensions to include texts, access to implementation, use of local and long haul communication optimization, integrity constraints must be compiled, forms-based user interface like tightly coupled, queries must be compiled, forms-based user environment, attention to autonomy	MRDSM	
DGs	Tightly coupled architecture, integrity constraints must be compiled, forms-based user interface	More like loosely coupled, integrity constraints must be compiled, forms-based user interface like tightly coupled, queries must be compiled, forms-based user environment, attention to autonomy	Tightly coupled, compiler, global query optimizer, auxiliary schema, IBM control, global optimization, extensions to include texts, access to implementation, use of local and long haul communication optimization, integrity constraints must be compiled, forms-based user interface like tightly coupled, queries must be compiled, forms-based user environment, attention to autonomy	OMNIBASE	
					Multibase
					Loosely coupled, DEC environment, query processing, distributed DBMS as a component DBMS

Table A.2. Significant Features

	Effort	Types of Components	CDM	DBMS	Reference
	Some Features				
ADDs	Extended E-R is used for integrity constraint enforcement at the federated schema level and at the external schema level. The primary CDM can be said to be relational since the internal command language is based on the relational algebra.	Relational	Relational	Relational	SCOOB [SPACCIPETRA et al. 1982]
CALIDA	Extended E-R is used for integrity constraint enforcement at the federated schema level and at the external schema level. The primary CDM can be said to be relational since the internal command language is based on the relational algebra.	Relational	Relational	Relational	SIRUS-DELTa [LITWIN et al. 1982]
DDTs	Relational	Relational	Relational	Relational	OMNIBASE [LITWIN et al. 1982]
DGs	Relational	Relational	Relational	Relational	Multibase [REINHOLD 1985]
					MRDSM [LITWIN 1985]
					OMNIBASE [LITWIN et al. 1982]
					SIRUS-DELTa [LITWIN et al. 1982]
					SCOOB [LITWIN et al. 1982]

Table A.1. Data Models of Various Mult-DBMS/FDBS Efforts

	Effort	Types of Components	CDM	DBMS	Reference
	Some Features				
Table A.1 summarizes the choice of a common data model and the types of components DBMSs integrated into some of the experiments partial prototyping systems. Table A.2 gives an empirical and partial list of significant features of some of these efforts. No effort is made to present all systems that exhibit FDBS features or to capture all important details of those mentioned.					
Table A.1 summarizes the choice of a common data model and the types of components DBMSs integrated into some of the experiments partial prototyping systems. Table A.2 gives an empirical and partial list of significant features of some of these efforts. No effort is made to present all systems that exhibit FDBS features or to capture all important details of those mentioned.					
Table A.1 summarizes the choice of a common data model and the types of components DBMSs integrated into some of the experiments partial prototyping systems. Table A.2 gives an empirical and partial list of significant features of some of these efforts. No effort is made to present all systems that exhibit FDBS features or to capture all important details of those mentioned.					

Appendix: Features of Some FDBS/Mult-DBMS Efforts

## What is semantics?

"the scientific study of the relations between signs and symbols and what they denote or mean".\*

So we have *sign and symbols* => model, language

- this is structured, limited, incomplete

what they *mean or denote* => real world information

- some times called real world semantics)
- this is unstructured, unlimited, possibly complete

One way to make a model useful is to talk about modeling in a specific *domain of discourse* [i.e., context]- hence limiting amount of information modeled (semantics captured).

## A critical issue in Information Modeling:

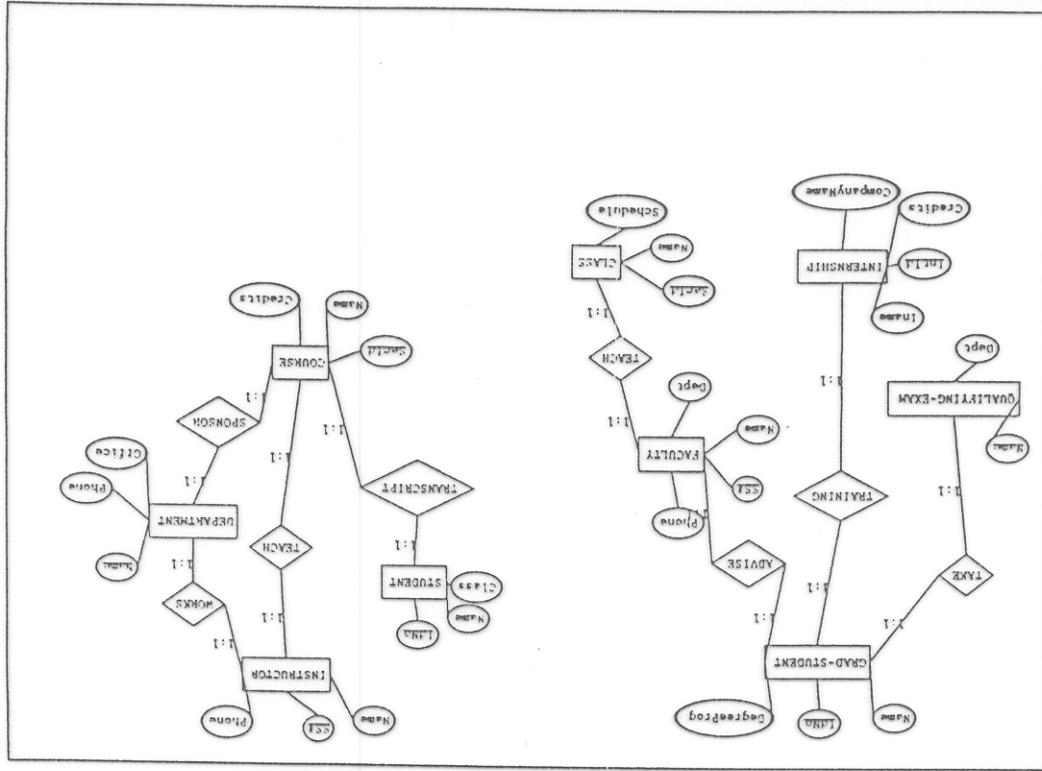
- Identify semantically related objects
- Determine schematic differences

### Modeling Semantics:

- Context
- Abstractions
- Domains
- States

\*Reference: William Woods, "What is a Link?."

[Sheth and Kashyap 92, papers in Sheth 91b, Kent 91]

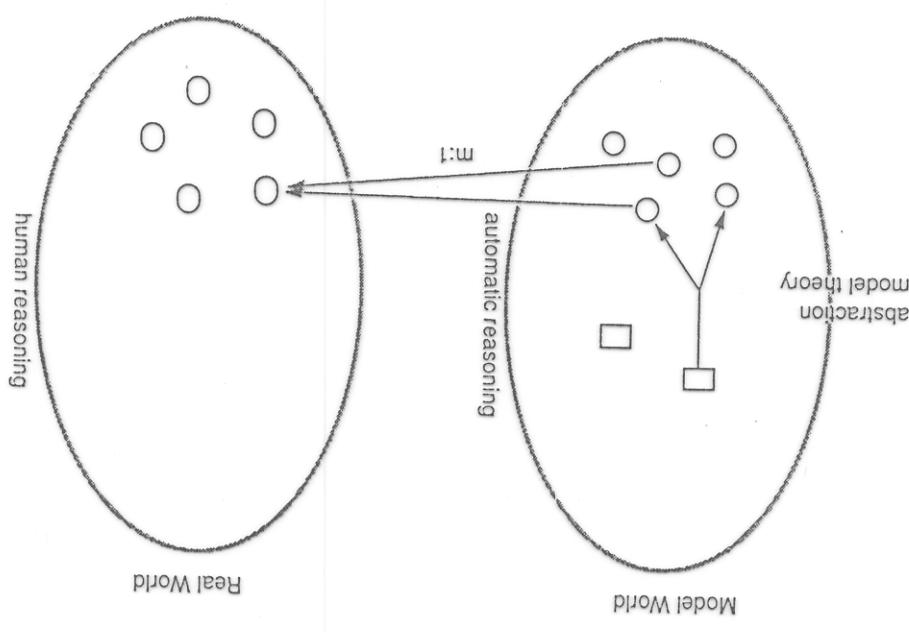


Schema S2

Schema S1

□ derived object

○ primitive object



Automatic vs. Human Reasoning