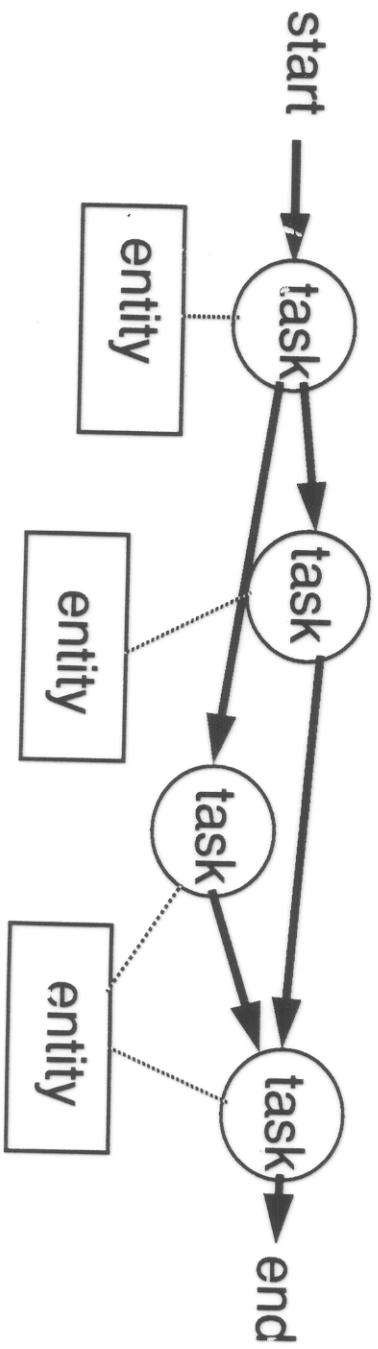


## Talk Outline

- Introduction to the basic concepts of workflow
- Formal model for specification and scheduling workflows
- Prototype and applications



# *Transactional Workflow Management*

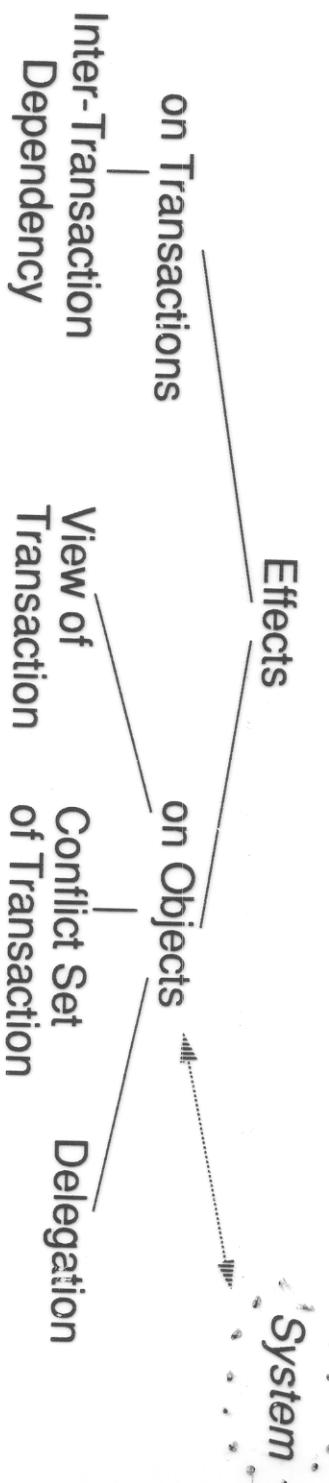
## Three Components:

- **Specification:**  
declarative, flexible specification of tasks/and dependencies  
extended transaction models [Elmagarmid 92] (e.g., Flexible Transaction [Elmagarmid/Rusinkiewicz et al. 90]; ACTA [Ramamirtham/ Chrysanthis 91/92], dependency specifications [Klein 91]
- **Scheduling:**  
efficient, maximal parallelism, exploit task and system semantics  
e.g., (L.O) [Cameron et al. 91]; VPL [Kuehn et al 92]
- **Executing:**  
manage execution of tasks/transactions on heterogeneous, autonomous component systems  
e.g., DOL System, Narada [Rusinkiewicz/UofH], Interbase [Elmagarmid/Purdue]

# *The ACTA Framework*

**Specification of the effects of transactions on other transactions and their effect on objects.**

- significant events associated with transactions
- histories and conditions on event occurrences
- objects, operations, and conflicts



## *Uses of ACTA*

- understand transactions properties
  - compare transaction models
- Not sufficient to generate schedule, or to support multiple transaction model interoperability.

## Beyond ACTA

- “postulative specifications” [Rusinkiewicz]
- “transaction specification toolkit”/ a-la-carte transaction
- specification of execution semantics [Sheth and Kalinichenko]

# Workflow Specification Model

- Significant events [ $e, e_i, e_j, \dots$ ]:  
events from tasks submitted by users/system (user events/  
system events); associate with a task agent *skeletons*
- Task Skeletons:  
an abstract representation of the actual task that hides  
irrelevant details of sequential computation of the task
- Intertask dependencies [ $D(e_1, \dots, e_n)$ ]:  
constraints on occurrence and temporal order of events
- System & application semantics

## *Workflow Specification Model*

- “Significant events” [ $e, e_i, e_j, \dots$ ]:  
events from tasks submitted by users/system (user events/  
system events); associate with a task agent
- Task Skeletons:
  - an abstract representation of the actual task that hides  
irrelevant details of sequential computation of the task
- Intertask dependencies [ $D(e_1, \dots, e_n)$ ]:  
constraints on occurrence and temporal order of events

# *Significant Events*

**Significant event types for database applications:** *st, ab, pr, cm*

**Possible attributes of an event type:**

- **Forcible:** the system can always force the execution
- **Rejectable:** the system can always reject the event
- **Delayable:** the system can delay execution of the event  
(every non-real-time significant events are delayable)

User Tasks	Forcible	Rejectable
cm	x	✓
ab	✗	x
pr	x	x
st	x	✗

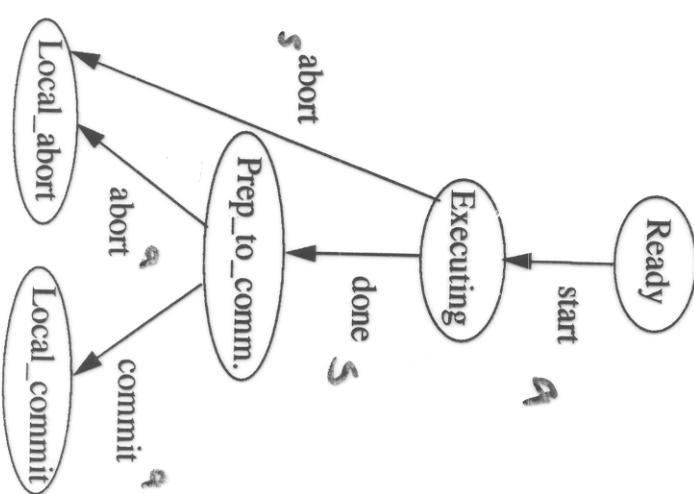
System Tasks	Forcible	Rejectable
cm	x	✓
ab	✓	x
pr	x	x
st	✓	✓

Table 1: Attribute Tables for Significant Events

## Task Skeleton

- different skeleton depending on application and the system
  - different states (e.g., no precommit)
  - different significant events submitted by application/user and system

Example:



## *Intertask Dependencies*

### Klein's primitives [KL91]:

- Order Dependency:  $e_1 < e_2$ .  
If both  $e_1$  and  $e_2$  occur, then  $e_1$  precedes  $e_2$ .  
Alternatively, in CTL: if  $e_2$  occurs,  $e_1$  cannot occur subsequently.  
Formally specified as:  $\text{AG}[e_2 \Rightarrow \text{AG} \sim e_1]$
- Existence Dependency:  $e_1 \rightarrow e_2$ .  
If event  $e_1$  occurs sometimes, then event  $e_2$  also occurs sometimes.

Alternatively, there is no computation such that  $e_2$  does not occur until a state  $s$  is reached where  $s$  satisfies [ $e_1$  is executed in  $s$ , and subsequently,  $e_2$  never occurs].  
Formally specified as:  $\sim E[\sim e_2 \cup (e_1 \wedge EG \sim e_2)]$

### – Conditional Existence Dependency [KL91]: $e_1 \rightarrow (e_2 \rightarrow e_3)$

### Examples from multidatabase transaction models:

- Commit Dependency [CR92]:  $cm_B < cm_A$
- Abort Dependency [CR92]:  $ab_B \rightarrow ab_A$

## Enforceable Dependencies

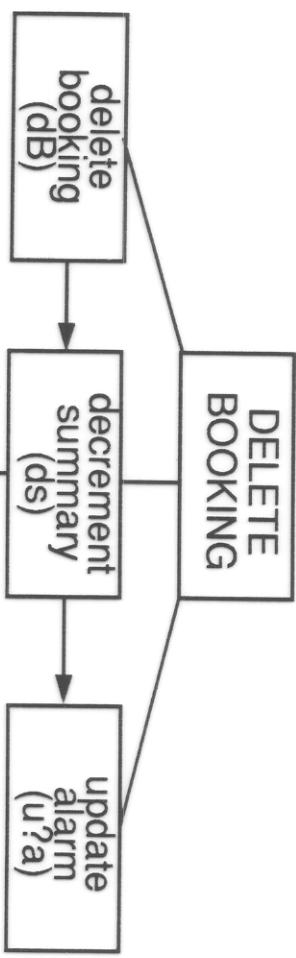
- Dependencies may not be enforceable.

For example,  $ab(A) \rightarrow cm(B)$

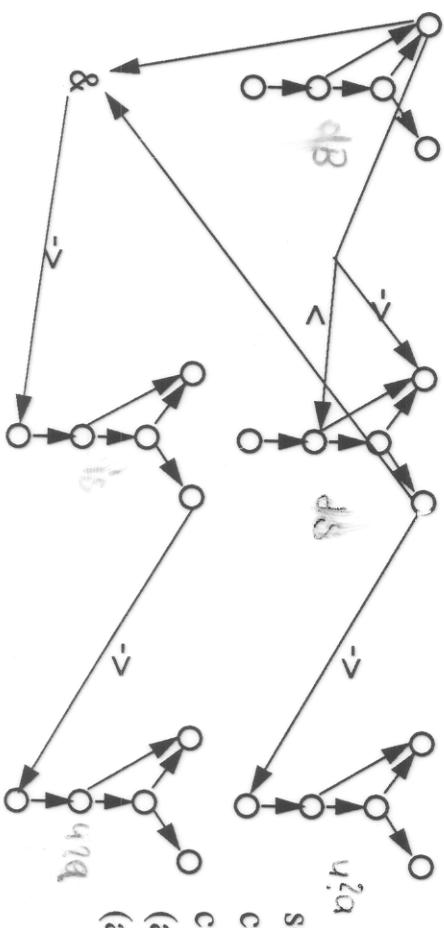
- Event attributes determine whether a dependency is enforceable. For example,
  - $e_1 \rightarrow e_2$  is run-time enforceable if
    - rejectable( $e_1$ )** [delay  $e_1$ ' until  $e_2$  is submitted, reject  $e_1$  if task 2 terminated without submitting  $e_2$ ],
    - or forcible( $e_2$ )** [force execution of  $e_2$  when  $e_1$  is accepted for execution].
  - $e_1 < e_2$  is run-time enforceable if
    - rejectable( $e_1$ )** [let  $e_2$  be executed when it is submitted, thereafter reject  $e_1$  if submitted],
    - or delayable( $e_2$ )** [delay  $e_2$  until either  $e_1$  has been accepted for execution, or task 1 has terminated without issuing  $e_1$ ].

# An Example

## Task Graph



## Intertask Dependencies



s(dB) -> s(ds)  
c(ds) -> s(u?a)  
c(is) -> s(u?a)  
(a(db) & c(ds)) -> s(is)  
(a(db) < d(ds)) -> a(ds)

## *Using CTL for dependency specification*

- formal semantics  
(propositional branching-time temporal logic: propositional logic and temporal operators)
- expressive, e.g., nesting of dependencies
- tools/algorithms for consistency and completeness checking
- limited temporal extension:  
dependencies that involve absolute clocks or relative-time service alarms (number of ticks), e.g., express: " $e_1 < e_2$  such that  $e_2$  occurs within t time units of  $e_1$ " or " $e_1 \rightarrow e_2$  such that  $e_2$  occurs no later than t time units after  $e_1$ "
- algorithms for automatic synthesis of automata for reactive systems  
(to develop scheduler to enforce intertask dependencies)

# *Using CTL for dependency specification*

- formal semantics  
(propositional branching-time temporal logic: propositional logic and temporal operators)
- expressive, e.g., nesting of dependencies
- tools/algorithms for consistency and completeness checking
- limited temporal extension:  
dependencies that involve absolute clocks or relative-time service alarms (number of ticks), e.g., express: " $e_1 <^+ e_2$  such that  $e_2$  occurs within  $t$  time units of  $e_1$ " or " $e_1 \rightarrow^+ e_2$  such that  $e_2$  occurs no later than  $t$  time units after  $e_1$ "
- algorithms for automatic synthesis of automata for reactive systems  
(to develop scheduler to enforce intertask dependencies)

## Dependency Automata

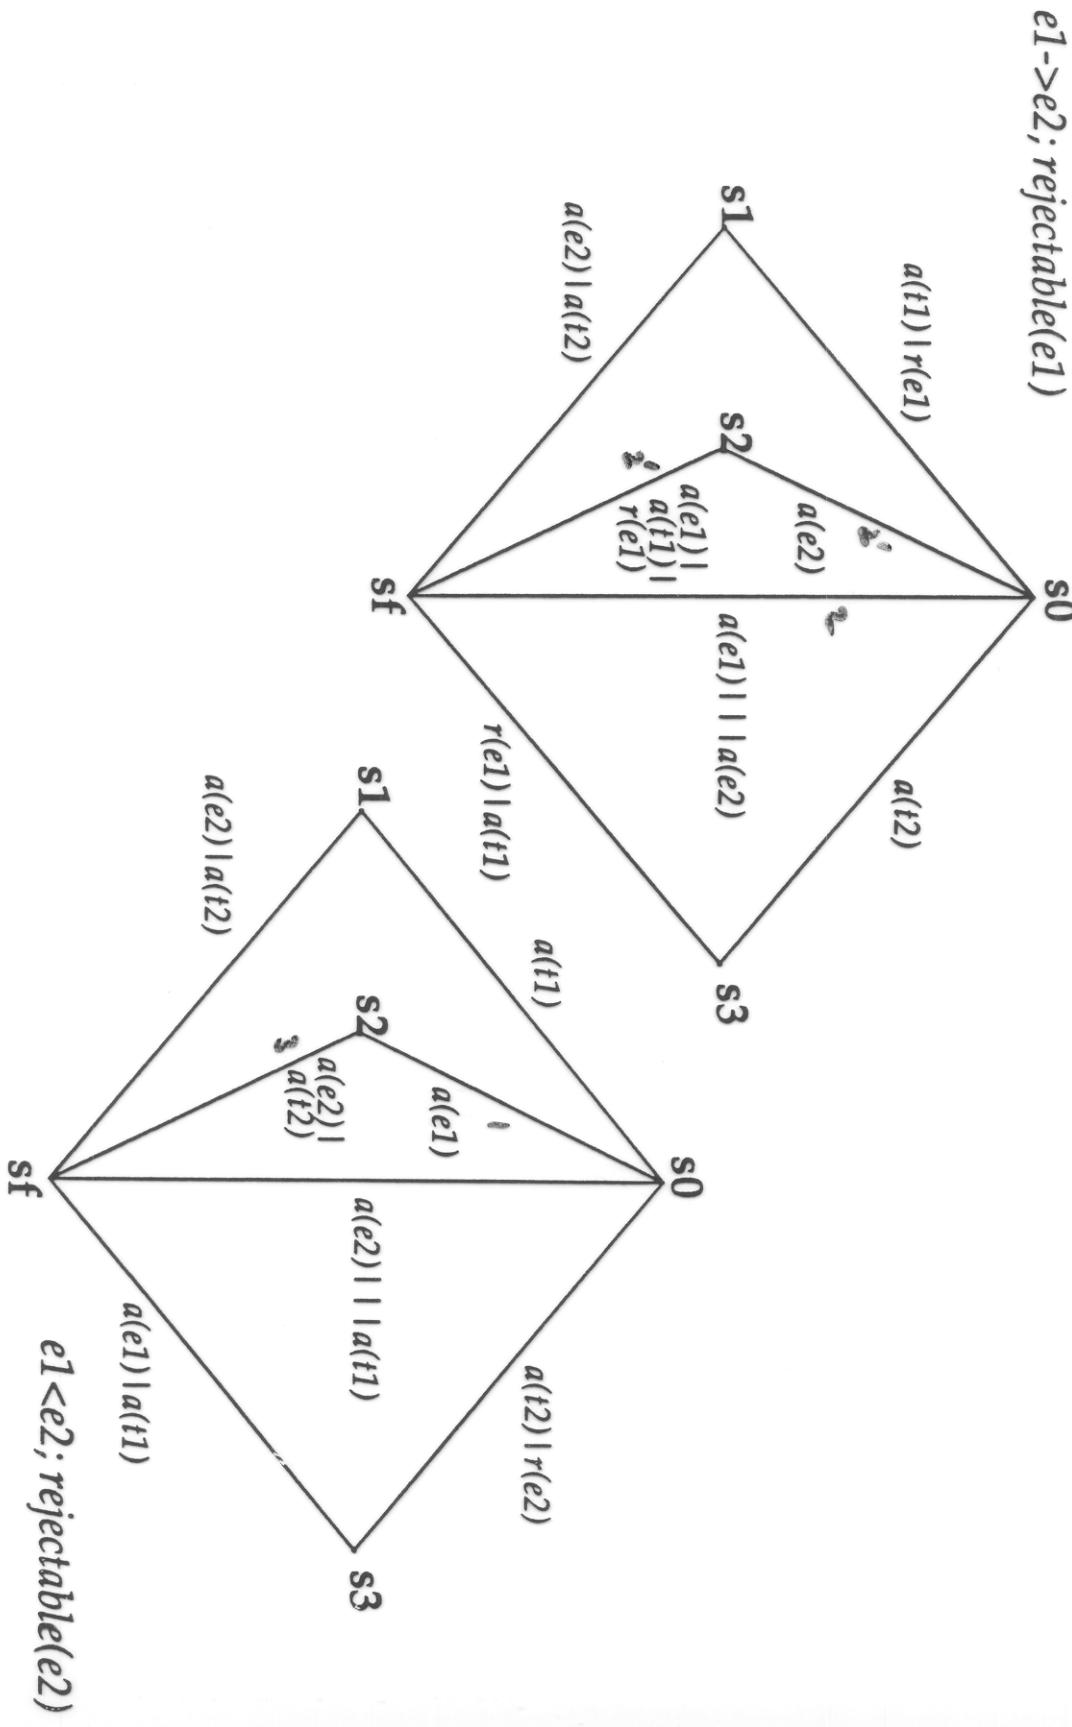
- Given  $D(e_1, \dots, e_k)$ , create a FSM  $A$  for enforcing  $D$ .
- Synthesized automatically from formula for  $D$ , and
- Properties of the events  $e_1, \dots, e_k$
- $\langle S, s', \Sigma, \rho \rangle$ , where
  - $s'$  is the initial state,  $\Sigma$  is the alphabet,
  - the transition relation.
- $s$  of  $\Sigma$  are of the form
  - $\epsilon$ , i.e.,  $A$  accepts  $\epsilon_1$  through  $\epsilon_m$
  - her relevant significant event  $e_j$  or its termination
- $m$ ), i.e.,  $A$  rejects  $e_1$  through  $e_m$ ,
- $\sigma_n$ , i.e., interleaving of accept operations in  $\sigma_1$  and  $\sigma_n$ , and
- i.e., accept op. in  $\sigma_i$  may occur before accept op. in

## Dependency Automata

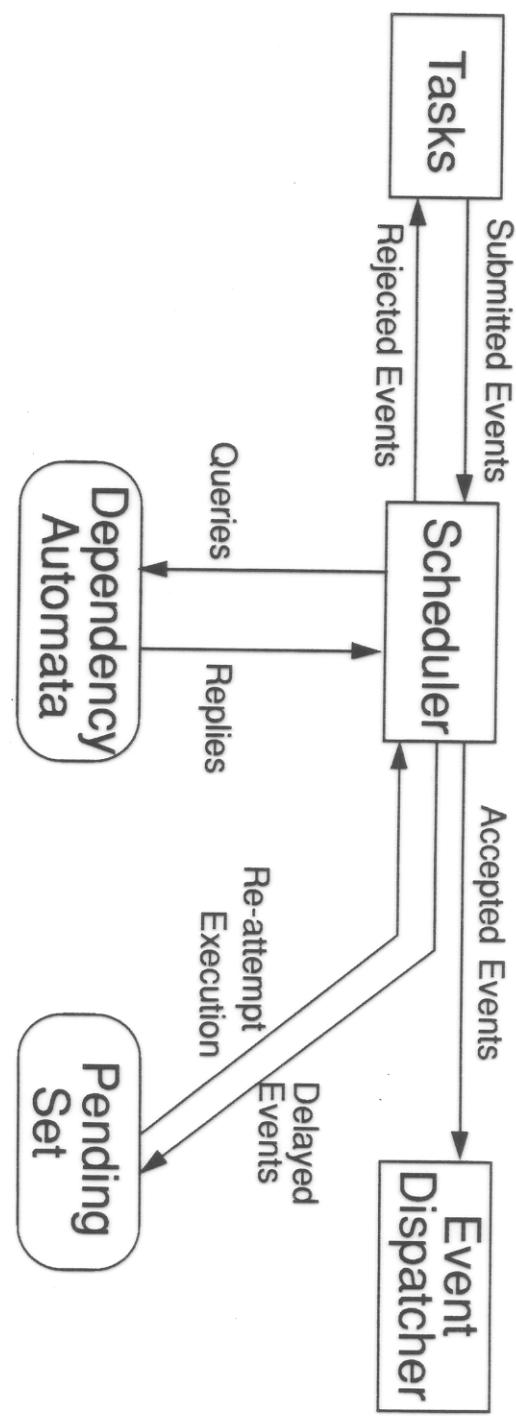
For a dependency  $D(e_1, \dots, e_k)$ , create a FSA  $A$  for enforcing  $D$ .

- $A$  can be synthesized automatically from
  - the CTL formula for  $D$ , and
  - the attributes of the events  $e_1, \dots, e_k$
- $A$  is a tuple  $\langle S, s', \Sigma, \rho \rangle$ , where
  - $S$  is a set of states,  $s'$  is the initial state,  $\Sigma$  is the alphabet, and  $\rho$  is the transition relation.
  - Elements of  $\Sigma$  are of the form $a(\varepsilon_1, \dots, \varepsilon_m)$ , i.e.,  $A$  accepts  $\varepsilon_1$  through  $\varepsilon_m$   
[ $\varepsilon_1$  is either relevant significant event  $e_1$  or its termination event  $t_1$ ,  
 $r(e_1, \dots, e_m)$ , i.e.,  $A$  rejects  $e_1$  through  $e_m$ ,  
 $\sigma_1 ||| \dots ||| \sigma_n$ , i.e., interleaving of accept operations in  $\sigma_1$  through  $\sigma_n$ , and  
 $\sigma_1, \dots, \sigma_n$ , i.e., accept op. in  $\sigma_i$  may occur before accept op. in  $\sigma_{i+1}$

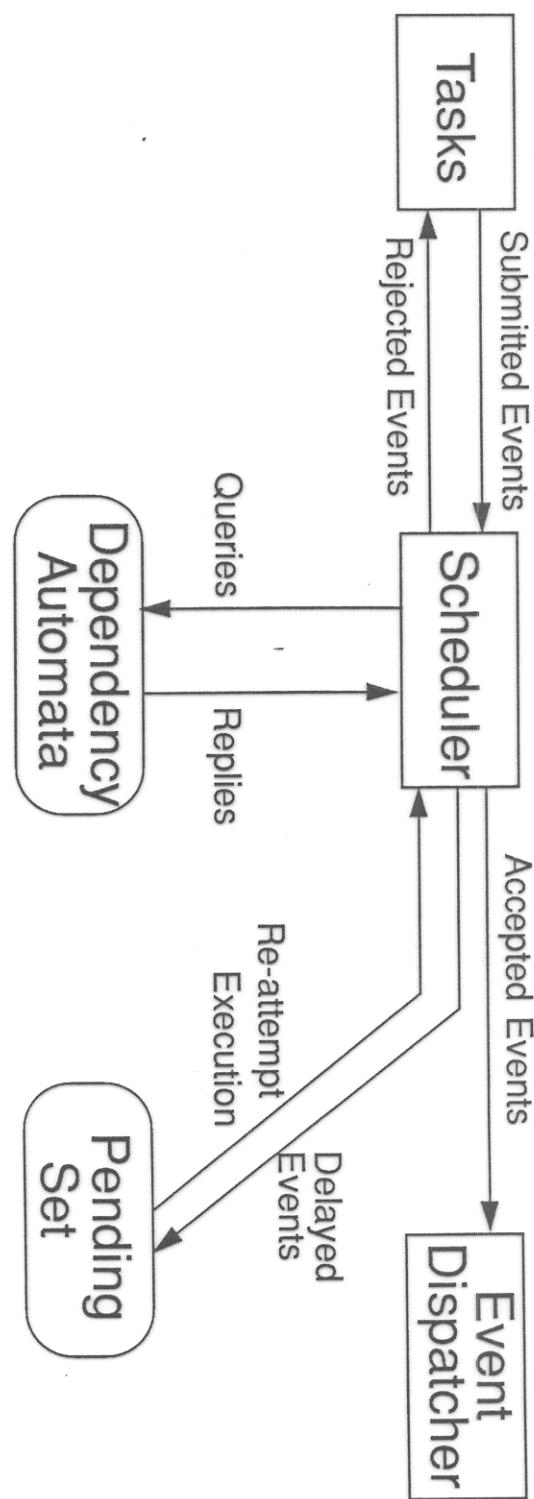
## Example Automata



# *Execution Model*



# Execution Model



## *forcing Multiple Dependencies*

the path corresponding to  $\epsilon$  from each relevant automaton.

pathset must:

$\epsilon$ ,

the current global state of the scheduler,

r-consistent,

all paths in the set must agree on the order of execution of each pair of

used or r-closed, and

any event that is accepted or rejected, path for each automaton referring to it must be included and must agree on whether to accept it or reject it

utable

accepted events must have been submitted and accepted events must be submitted or be forcible.

# *Enforcing Multiple Dependencies*

Pathset: one path corresponding to  $\epsilon$  from each relevant dependency automaton.

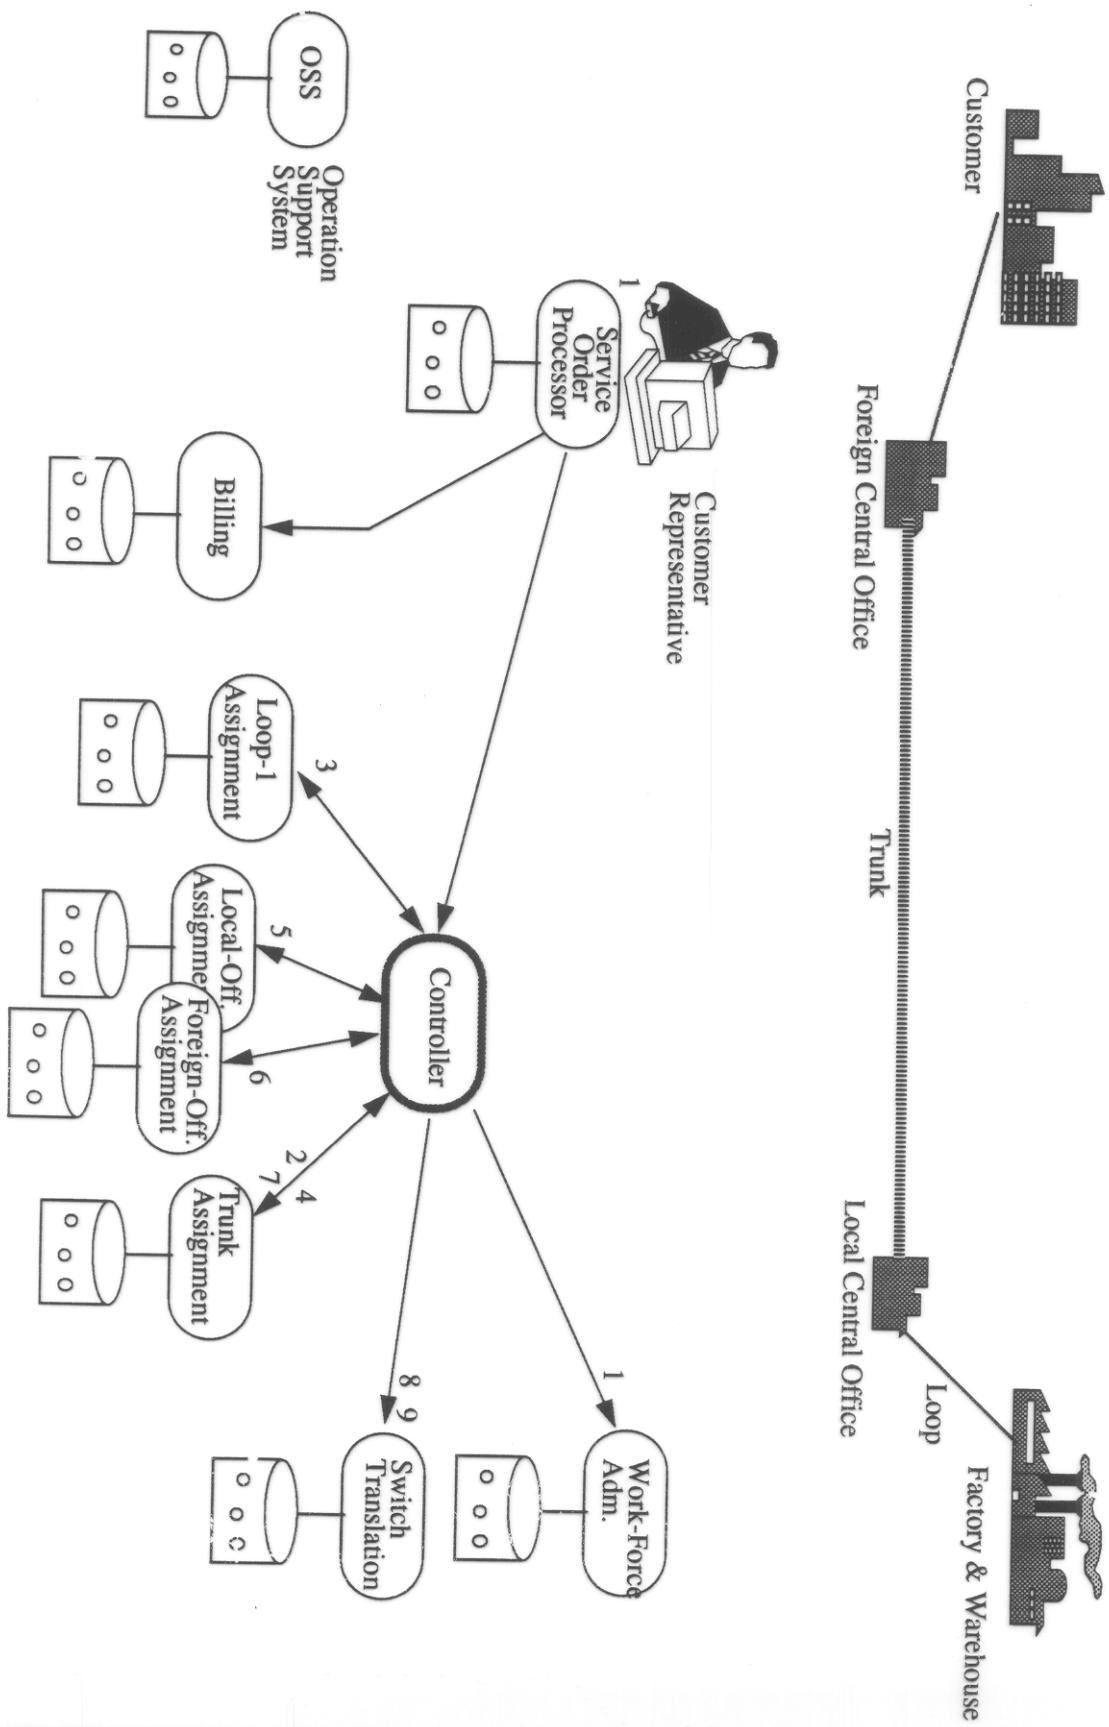
A *desired pathset* must:

- accept  $\epsilon$ ,
- begin in the current global state of the scheduler,
- be order-consistent,  
i.e., different paths in the set must agree on the order of execution of each pair of events
- be a-closed or r-closed, and
  - i.e., for any event that is accepted or rejected, path from each automaton referring to that event must be included and must agree on whether to accept it or reject it
- be executable
  - i.e., all rejected events must have been submitted and all accepted events must have been submitted or be forcible.

## Scheduler Operation (An Example)

- Consider only  $e_1 < e_2$  and  $e_1 \rightarrow e_2$  dependencies, where both  $e_1$  and  $e_2$  are rejectable (e.g., all dependencies for SAGAs can be expressed using these). Corresponding automata  $A_<$  and  $A_{\rightarrow}$ .
- $e_1$  is submitted.
  - $a(e_1)$  in  $A_<$ . No path in  $A_{\rightarrow}$  with  $e_1$ .  $e_1$  added to pending set.
  - $e_2$  is submitted.
  - $A_{\rightarrow} : a(e_2); a(e_1)$  and  $a(e_2) \parallel a(e_1)$ .
  - $a$ -closure forces searching  $A_<$  for a path that accepts both  $e_1$  and  $e_2$ . Only such path is  $a(e_1); a(e_2)$  which is not order-consistent with  $a(e_2); a(e_1)$ .
  - Viable pathset is  $\{a(e_1); a(e_2), a(e_2) \parallel a(e_1)\}$ .
  - Partial order consistent with this is  $e_1$ , and then  $e_2$ .

# An Example Workflow in a Multisystem Application: Provisioning a Telecommunication Service



## *About the environment*

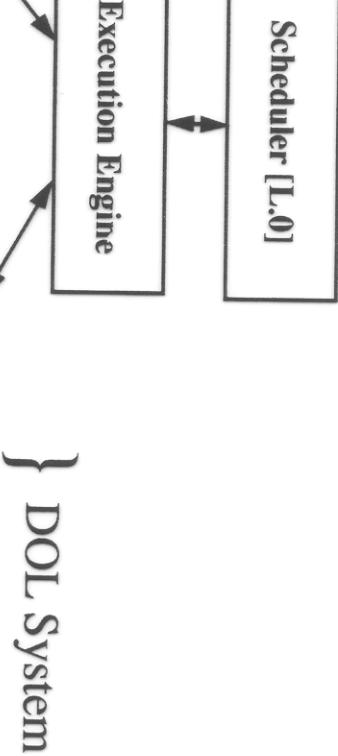
- multiple existing heterogeneous “closed” application systems
- each system developed independently to automate a business function
- each with own databases multiple existing heterogeneous “closed” application systems
- each has predefined interface (“contracts”)
- multisystem application implemented using dedicated controller
  - hard-coded, difficult to change work-flow
  - use of queued message paradigm
  - no use of transaction paradigm for multisystem application
    - application specific and application managed concurrency control and recovery

Timestamp

# PROMT

(A system for Processing Multidatabase Transactions)

Flexible Transaction  
↓



OSSs (application systems) emulated using DBMSs.

[Bellcore, UofH]

## *Beyond PROMT*

- Scheduler based on a formal framework
- Concurrency Control and Recovery that exploit application and system semantics (implementation in progress)
- Interface with real OSes
- Graphical specification and testing of workflows (in progress)
- Graphical instantiation and monitoring of workflows
- Detailed specifications of individual tasks (messages to OSSs)
- Data dependencies and translations using tables, rules, constants, and syntactic changes

## *Conclusions*

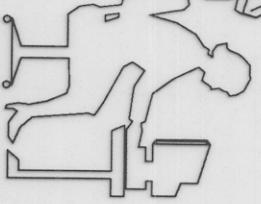
- Gained detailed understanding on issues of transaction workflows.
- Studied and demonstrated applicability of relaxed multidatabase transaction.
- **Developed formal approach to specifying and executing workflows.**
- Completed large-scale prototype.  
Demonstrated with a real application.  
Deployment considered.

## *PROMT Conclusions*

- Flexible Transaction model seems natural to model multi-system work-flows
  - Some extensions/improvements performed/desired:
    - specification of acceptable states
    - additional execution dependencies
- Use of relaxed transaction technology appropriate
  - already practiced informally, but relaxed transactions provide better and precise understanding of relaxation of ACID properties
  - separation of control from individual (sub) transactions
  - functionally better than current practice of message queues or store-and-forward with manual intervention
  - better modeling of component systems (subtransaction state transitions)
- Yet to understand how to handle manual interventions

## *Prototype 2: SPIRIT*

Recent Research



Service  
Order  
Representative  
- enter new orders  
- monitor progress

Service  
Developer  
and  
System  
Analyst

- specify needs
- test/simulate

Allayst

**Information Specification & Modeling**

Workflow Control System

**Scherr modeling, analysis, and integration in automating integration processes,** Monitoring Autom. Engg.

## - semantic and schematic taxonomy, context representation

- Information integrity/~~consistency~~

– multidatabase transaction dependency constraints, Polytractions

Complex applications [multimedia, personal intelligent communication]  
OSS  
Network

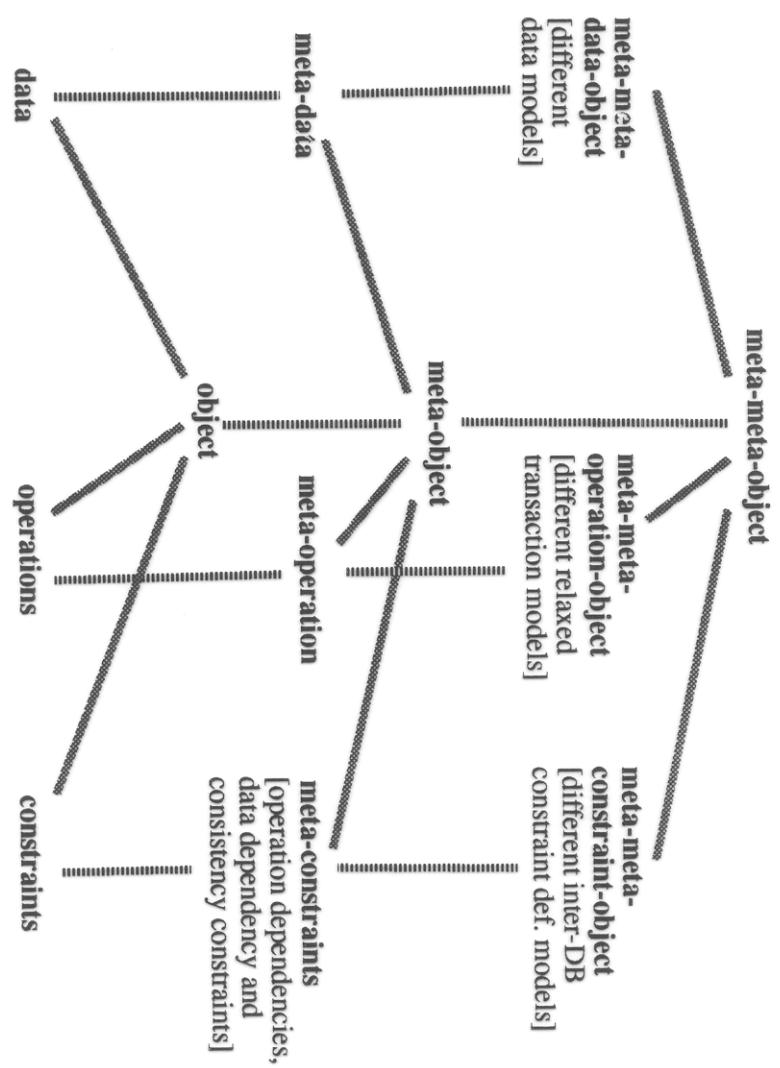
- controlling cooperative agent, [information brokering]

Existing and New Systems

OSS: Operation Support System

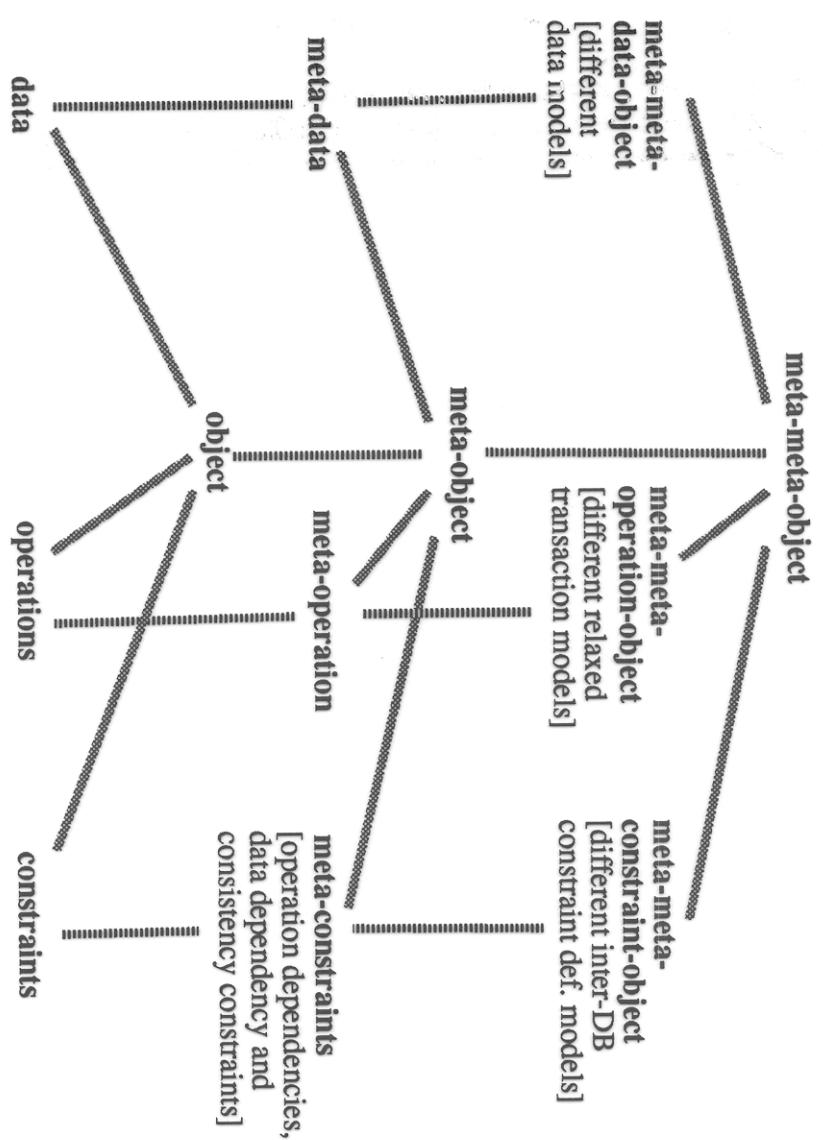
[MCC, Bellcore, UofH, Ameritech]

# Modeling, Organizing and Managing Information Systems



[Sheth, Kalinichenko 92]

# Modeling, Organizing and Managing Information Systems



[Sheth, Kalinichenko 92]

# *Issues in Interoperability in Information Systems*

