


Talk Outline

- Basic concepts and specification of workflows
- Formal model for specification and scheduling of intertask dependencies
- Prototypes and applications
- Work in progress



What is a (transactional) workflow?

An activity that involves coordinated execution of multiple related tasks by different entities.

Specification issues:

- Individual Tasks:
 - task format: message, contract, form *transaction*
 - task structure: externally visible states of the task, initial state, termination states, significant events and *their attributes*
 - task (operation) semantics, e.g., compatibility, relaxed isolation
- Individual Entities:
 - type of entity: human, application system, DBMS
 - system properties/semantics, e.g., isolation granularity, order preservation, idempotency, monotonicity

What is a (transactional) workflow?

- Task Coordination requirements:
 - intertask dependencies and data exchange
- Intra- and inter-workflow Execution requirements:
 - failure atomicity (A)
 - execution atomicity (I)
 - workflow recovery
 - inter-workflow concurrency

Workflow Examples

Environment	Application
office computing	mail routing loan processing meeting scheduling course organizing
data processing	processing a purchase order
manufacturing	product life-cycle
telecommunication	establishing or changing a service/circuit

Closely related terms/issues:

Multi-system applications [Bellcore/UofH], task flow [Dayal], long-running activities [DEC], application multi-activities [Kalinechenko], extended transaction models [Elmagarmid book], third generation TP monitor [SIGMOD93]

Related research areas [different types of tasks, different types of entities]: cooperative activity [Bellcore,..], collaborative distributed problem solving [UFL,..], DAI [DAKE, MCC,..], learning, self-adapting software agents [CMU,..]

Transactional Workflow Management

Three Components:

- **Specification:**
declarative, flexible specification of tasks, dependencies, execution requirements
extended transaction models [Elmagarmid 92] (e.g., Flexible Transaction [Elmagarmid/Rusinkiewicz et al. 90]; ACTA [Ramamritham/Chrysanthis 91/92], dependency specifications [Klein 91])
- **Scheduling:**
efficient, maximal parallelism, exploit task and system semantics
e.g., (L.O) [Cameron et al. 91]; VPL [Kuehn et al 92]
- **Executing:**
manage execution of tasks/transactions on heterogeneous, autonomous component systems
e.g., DOL System, Narada [Rusinkiewicz/UofH], Interbase/RSI [Elmagarmid/Purdue], ESS [MCC/Car-not]

Transaction Models and Specifications

- ACID transactions and their nested derivatives
Problems: inflexible, difficult to implement in multi-systems.
- Queued message systems and “chaining of transactions”.
Problems: insufficient control over transaction properties, interactions among concurrent activities difficult.
- Extended/Relaxed Transaction Models:
 - Sagas and Nested Sagas [Garcia-Molina et al. 88, 90]
 - ConTracts [Reuter 89]
 - Flexible Transactions [Elmagarmid et al 90, Rusinkiewicz et al 90]
 - Multi-transaction Activities [Garcia-Molina et al. 90]
 - Long-Running Activities [Dayal et al. 91]
 - Relaxed transactions in Carnot [Cannata 91]
 - The DOM project [Buchmann et al 92]
 - Open Nested Transactions [Weikum & Schek 92] and Others (e.g., in [Elmagarmid 92], SIGMOD93)

Significant Events

Significant event types for database applications: *st*, *ab*, *pr*, *cm*

Possible attributes of an event type:

- Forcible: the system can always force the execution
- Rejectable: the system can always reject the event
- Delayable: the system can delay execution of the event
(every non-real-time significant events are delayable)




Event	Forcible?	Rejectable?	Delayable?
cm	N	Y	Y
ab	Y	N	N
pr	N	N	N
st	Y	Y	Y

Usual attribute assignments for transactions
in database applications and DBMSs

Task Skeleton

- different skeleton depending on application and the system
 - different states (e.g., no precommit)
 - different significant events submitted by application/user and system

Examples:



Intertask Dependencies

Preconditions for initiating each scheduler-controllable transition in a task.

Klein's primitives [KL91]:

- Order Dependency: $e_1 < e_2$.

If both e_1 and e_2 occur, then e_1 precedes e_2 .

Alternatively, in CTL: if e_2 occurs, e_1 cannot occur subsequently.

Formally specified as: $AG[e_2 \Rightarrow AG \sim e_1]$

- Existence Dependency: $e_1 \rightarrow e_2$.

If event e_1 occurs sometimes, then event e_2 also occurs sometimes.

Alternatively, there is no computation such that e_2 does not occur until a state s is reached where s satisfies $[e_1]$ is executed in s , and subsequently, e_2 never occurs].

Formally specified as: $\sim E[\sim e_2 \cup (e_1 \wedge EG \sim e_2)]$

- Conditional Existence Dependency [KL91]: $e_1 \rightarrow (e_2 \rightarrow e_3)$

Examples from multidatabase transaction models:

- Commit Dependency [CR92]: $cm_B < cm_A$
- Abort Dependency [CR92]: $ab_B \rightarrow ab_A$

Enforceable Dependencies

- Dependencies may not be enforceable.

For example, $ab(A) \rightarrow cm(B)$

- Event attributes determine whether a dependency is enforceable. For example,
 - $e_1 \rightarrow e_2$ is run-time enforceable if
rejectable(e_1) [delay e_1 until e_2 is submitted, reject e_1 if task 2 terminated without submitting e_2],
or forcible(e_2) [force execution of e_2 when e_1 is accepted for execution].
 - $e_1 < e_2$ is run-time enforceable if
rejectable(e_1) [let e_2 be executed when it is submitted, thereafter reject e_1 if submitted],
or delayable(e_2) [delay e_2 until either e_1 has been accepted for execution, or task 1 has terminated without issuing e_1].

Beyond Dependencies -- Task Coordination Requirements

Statically -- a precondition for starting a task or initiating a transition in a task.

Preconditions may be specified with dependencies involving:

- execution states of other tasks
- output values of other tasks
- external variables (events outside the workflow, time,..)

E.g., execution dependencies, data/value dependencies, temporal dependencies in Flexible Transactions [Elmagarmid et al 90], ConTracts [Reuter 89], Multitransactions [Garcia-Molina et al 90], Multidatabase Transactions [Rusinkiewicz et al 92]....

Dynamically--

Created when executing a workflow.

Long-running activities [Dayal et al 91], Polytransactions [Rusinkiewicz and Sheth 91].

Scheduler Operation (An Example)

Consider only $e_1 < e_2$ and $e_1 \rightarrow e_2$ dependencies, where both e_1 and e_2 are rejectable (e.g., all dependencies for SAGAs can be expressed using these). Corresponding automata $A_{<}$ and A_{\rightarrow} .

- e_1 is submitted.
 - $a(e_1)$ in $A_{<}$. No path in A_{\rightarrow} with e_1 . e_1 added to pending set.
- e_2 is submitted.
 - $A \rightarrow: a(e_2); a(e_1)$ and $a(e_2) ||| a(e_1)$.
 - a -closure forces searching $A_{<}$ for a path that accepts both e_1 and e_2 . Only such path is $a(e_1); a(e_2)$ which is not order-consistent with $a(e_2); a(e_1)$.
 - Viable pathset is $\{a(e_1); a(e_2), a(e_2) ||| a(e_1)\}$.
 - Partial order consistent with this is e_1 and then e_2 .

ion/Task and System Semantics

Task, System	Impact (CCon. Control, R: Recovery) ↓
appl)	fewer exclusive locks (CC)
	no global commitment (CC)
porousness (sys)	early release of locks (CC)
	resubmit transactions (R)
	roll-forward recovery (R)

Conclusions

- Gained detailed understanding on issues of transaction workflows.
- Studied and demonstrated applicability of relaxed multidatabase transaction.
- Developing a generic model workflow.
- **Developed formal approach to specifying and executing workflows.**
- Completed large-scale prototype.
Demonstrated with a real application.
Deployment being considered.

External Publications

E. Jane Cameron, Linda Ness, and Amit Sheth, "A Universal Executor for Flexible Transactions which Achieves Maximal Parallelism," in the *Proc. of the 1st Intl. Workshop on Interoperability in Multidatabase Systems* (IMS '91), April 1991.

Authors1, "Executing Multidatabase Transactions," in the *Proc. of the 25th Intl. Conf. on Systems Sciences*, January 92.

Authors1, "Using Flexible Transactions to Support Multi-system Applications", presented at the *1992 Workshop on Heterogeneous Databases and Semantic Interoperability*, U S West, Boulder, February 1992.

Authors1, "Applying Multidatabase Transactions for Service Order Processing", in the *Proc. of the 18th Intl. Conf. on Very Large Databases* (VLDB'92), August 1992.

Authors2, "Concurrency Control and Recovery of Multidatabase Work Flows in Telecommunication Applications. in the *Proc. of ACM SIGMOD Conference*, May 1993.

D. Woelk, P. Attie, P. Cannata, G. Meredith, A. Singh, M. Singh, and C. Tomlinson, "Task Scheduling using Intertask Dependencies in Carnot," in the *Proc. of ACM SIGMOD Conference*, May 1993.

P. Attie, M. Singh, M. Rusinkiewicz, and A. Sheth, "Specifying and Enforcing Intertask Dependencies," MCC Technical Report # Carnot-245-92, November 1992. Abridged version in the *Proc. of the 19th Intl Conf. on Very Large Data Bases*, August 1993.

A. Sheth and M. Rusinkiewicz, "On Transactional Workflows," in *Data Engineering Bulletin*, 16 (2), June 1993.

Authors1 (different orders): M. Ansari, L. Ness, M. Rusinkiewicz, A. Sheth
Authors2: W. Jin, L. Ness, M. Rusinkiewicz, A. Sheth